

# texas\_agriculture\_eda

November 14, 2025

## 1 Texas Agriculture & Climate Data - EDA

### 1.1 Exploring patterns in crop yields and weather (2000-2023)

**Team:** Carter Dobbs, Johann Steinhoff, Jay Suh  
**Course:** CS 4347 - Introduction to Machine Learning  
**November 2025**

---

This notebook explores our merged dataset combining USDA agricultural data with NOAA climate records for Texas. We're interested in understanding how weather patterns might affect crop production, which will inform our predictive modeling approach.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

### 1.2 Loading the data

```
[2]: # Load our merged dataset
df = pd.read_csv("./texas_agriculture_with_climate_2000_2023.csv", ↴
    low_memory=False)

print(f"Dataset dimensions: {df.shape[0]} rows x {df.shape[1]} columns")
print(f"Total data points: {df.shape[0] * df.shape[1]} points")
```

Dataset dimensions: 398,204 rows x 120 columns  
Total data points: 47,784,480 points

### 1.3 Understanding our features

The 120 columns break down into a few categories:

- Agricultural metadata (crop types, measurement units, etc.)
- Climate measurements (monthly temp/precip data)
- Features we engineered (seasonal and annual aggregates)

```
[3]: # Group columns by category for easier analysis later
usda_columns = [
    "SOURCE_DESC",
```

```

    "SECTOR_DESC",
    "GROUP_DESC",
    "COMMODITY_DESC",
    "CLASS_DESC",
    "PRODN_PRACTICE_DESC",
    "UTIL_PRACTICE_DESC",
    "STATISTICCAT_DESC",
    "UNIT_DESC",
    "SHORT_DESC",
    "DOMAIN_DESC",
    "DOMAINCAT_DESC",
    "AGG_LEVEL_DESC",
]
location_columns = [
    "STATE_ANSI",
    "STATE_FIPS_CODE",
    "STATE_ALPHA",
    "STATE_NAME",
    "ASD_CODE",
    "ASD_DESC",
    "COUNTY_ANSI",
    "COUNTY_CODE",
    "COUNTY_NAME",
    "REGION_DESC",
    "ZIP_5",
    "WATERSHED_CODE",
    "WATERSHED_DESC",
    "CONGR_DISTRICT_CODE",
    "COUNTRY_CODE",
    "COUNTRY_NAME",
    "LOCATION_DESC",
    "COUNTY_FIPS",
]
temporal_columns = [
    "YEAR",
    "FREQ_DESC",
    "BEGIN_CODE",
    "END_CODE",
    "REFERENCE_PERIOD_DESC",
    "WEEK_ENDING",
    "LOAD_TIME",
]
# Monthly climate data
months = [

```

```

    "JAN",
    "FEB",
    "MAR",
    "APR",
    "MAY",
    "JUN",
    "JUL",
    "AUG",
    "SEP",
    "OCT",
    "NOV",
    "DEC",
]
precip_columns = [f"PRECIP_{m}" for m in months]
tmax_columns = [f"TMAX_{m}" for m in months]
tmin_columns = [f"TMIN_{m}" for m in months]
tavg_columns = [f"TAVG_{m}" for m in months]
cdd_columns = [f"CDD_{m}" for m in months] # cooling degree days
hdd_columns = [f"HDD_{m}" for m in months] # heating degree days

# Aggregated climate features we created
engineered_columns = [
    "GROWING_SEASON_PRECIP",
    "GROWING_SEASON_TEMP_AVG",
    "GROWING_SEASON_TEMP_MAX",
    "GROWING_SEASON_TEMP_MIN",
    "ANNUAL_PRECIP",
    "ANNUAL_TEMP_AVG",
    "ANNUAL_CDD",
    "ANNUAL_HDD",
]

```

## 1.4 Column breakdown

Let's see what we're actually working with:

```
[4]: # Show what columns we have in each category
print("USDA AGRICULTURAL COLUMNS:")
print(f"({len([c for c in usda_columns if c in df.columns])} columns)")
for col in usda_columns:
    if col in df.columns:
        print(f" - {col}")

print("\n\nLOCATION COLUMNS:")
print(f"({len([c for c in location_columns if c in df.columns])} columns)")
for col in location_columns:
    if col in df.columns:

```

```

        print(f" - {col}")

print(f"\n\nTIME COLUMNS:")
print(f"({len([c for c in temporal_columns if c in df.columns])} columns)")
for col in temporal_columns:
    if col in df.columns:
        print(f" - {col}")

print(f"\n\nMONTHLY CLIMATE COLUMNS:")
climate_cols = (
    precip_columns
    + tmax_columns
    + tmin_columns
    + tavg_columns
    + cdd_columns
    + hdd_columns
)
monthly_count = len([c for c in climate_cols if c in df.columns])
print(f"{monthly_count} columns")
print(" Sample columns:")
print(f"    Precipitation: {[c for c in precip_columns if c in df.columns][:3]} ...")
print(f"    Max Temp: {[c for c in tmax_columns if c in df.columns][:3]} ...")
print(f"    Min Temp: {[c for c in tmin_columns if c in df.columns][:3]} ...")
print(f"    Avg Temp: {[c for c in tavg_columns if c in df.columns][:3]} ...")
print(f"    Cooling Degree Days: {[c for c in cdd_columns if c in df.columns][:3]} ...")
print(f"    Heating Degree Days: {[c for c in hdd_columns if c in df.columns][:3]} ...")

print(f"\n\nENGINEERED CLIMATE FEATURES:")
print(f"({len([c for c in engineered_columns if c in df.columns])} columns)")
for col in engineered_columns:
    if col in df.columns:
        print(f" - {col}")

```

USDA AGRICULTURAL COLUMNS:

- (13 columns)
- SOURCE\_DESC
- SECTOR\_DESC
- GROUP\_DESC
- COMMODITY\_DESC
- CLASS\_DESC
- PRODN\_PRACTICE\_DESC
- UTIL\_PRACTICE\_DESC
- STATISTICCAT\_DESC
- UNIT\_DESC

- SHORT\_DESC
- DOMAIN\_DESC
- DOMAINCAT\_DESC
- AGG\_LEVEL\_DESC

**LOCATION COLUMNS:**

- (18 columns)
- STATE\_ANSI
  - STATE\_FIPS\_CODE
  - STATE\_ALPHA
  - STATE\_NAME
  - ASD\_CODE
  - ASD\_DESC
  - COUNTY\_ANSI
  - COUNTY\_CODE
  - COUNTY\_NAME
  - REGION\_DESC
  - ZIP\_5
  - WATERSHED\_CODE
  - WATERSHED\_DESC
  - CONGR\_DISTRICT\_CODE
  - COUNTRY\_CODE
  - COUNTRY\_NAME
  - LOCATION\_DESC
  - COUNTY\_FIPS

**TIME COLUMNS:**

- (7 columns)
- YEAR
  - FREQ\_DESC
  - BEGIN\_CODE
  - END\_CODE
  - REFERENCE\_PERIOD\_DESC
  - WEEK\_ENDING
  - LOAD\_TIME

**MONTHLY CLIMATE COLUMNS:**

(72 columns)

**Sample columns:**

```
Precipitation: ['PRECIP_JAN', 'PRECIP_FEB', 'PRECIP_MAR'] ...
Max Temp: ['TMAX_JAN', 'TMAX_FEB', 'TMAX_MAR'] ...
Min Temp: ['TMIN_JAN', 'TMIN_FEB', 'TMIN_MAR'] ...
Avg Temp: ['TAVG_JAN', 'TAVG_FEB', 'TAVG_MAR'] ...
Cooling Degree Days: ['CDD_JAN', 'CDD_FEB', 'CDD_MAR'] ...
Heating Degree Days: ['HDD_JAN', 'HDD_FEB', 'HDD_MAR'] ...
```

ENGINEERED CLIMATE FEATURES:

- (8 columns)
- GROWING\_SEASON\_PRECIP
- GROWING\_SEASON\_TEMP\_AVG
- GROWING\_SEASON\_TEMP\_MAX
- GROWING\_SEASON\_TEMP\_MIN
- ANNUAL\_PRECIP
- ANNUAL\_TEMP\_AVG
- ANNUAL\_CDD
- ANNUAL\_HDD

## 1.5 Missing data check

```
[5]: # Check which columns have missing values
missing_data = pd.DataFrame(
    {
        "Column": df.columns,
        "Missing": df.isnull().sum().values,
        "Percent": (df.isnull().sum().values / len(df) * 100),
    }
)
missing_data = missing_data[missing_data["Missing"] > 0].sort_values(
    "Percent", ascending=False
)

print(f"Columns with missing data: {len(missing_data)}/{len(df.columns)}")
print(f"Total missing: {missing_data['Missing'].sum():,}")
print(
    f"Overall completeness: {100 - (missing_data['Missing'].sum() / (len(df) * len(df.columns)) * 100):.1f}%""
)
print(f"\nTop columns with missing values:")
print(missing_data.head(10).to_string(index=False))
```

Columns with missing data: 88/120

Total missing: 3,348,024

Overall completeness: 93.0%

Top columns with missing values:

Column	Missing	Percent
ZIP_5	398204	100.000000
REGION_DESC	398204	100.000000
CONGR_DISTRICT_CODE	398204	100.000000
WATERSHED_DESC	398204	100.000000
WEEK_ENDING	398204	100.000000
CV_%	211970	53.231509

COUNTY_ANSI	14133	3.549186
PRECIP_JAN	14133	3.549186
PRECIP_SEP	14133	3.549186
PRECIP_FEB	14133	3.549186

### 1.5.1 Column variance analysis

```
[6]: # Compute variance for all numeric columns to identify high-variance features
num_cols = df.select_dtypes(include=[np.number]).columns.tolist()
desc = df[num_cols].describe().T
desc["median"] = df[num_cols].median()
desc["skew"] = df[num_cols].skew()

cols_to_show = ["count", "min", "median", "mean", "std", "max", "skew"]
display(desc[cols_to_show].round(3))

top_var = df[num_cols].var().sort_values(ascending=False).head(10)
print("\nTop 10 numeric columns by variance:")
display(top_var)

print("\nNote: High variance in VALUE is expected since it measures different")
print(
    "units (dollars, bushels, acres). Feature scaling will be important for"
    "modeling."
)
```

	count	min	median	mean	std	\
STATE_ANSI	398204.0	48.000	48.000	48.000	0.000	
STATE_FIPS_CODE	398204.0	48.000	48.000	48.000	0.000	
ASD_CODE	398204.0	11.000	51.000	49.134	27.868	
COUNTY_ANSI	384071.0	1.000	233.000	246.936	147.435	
COUNTY_CODE	398204.0	1.000	245.000	273.592	200.688	
...	...	...	...	...	...	
GROWING_SEASON_TEMP_MIN	384071.0	52.683	66.467	65.701	4.260	
ANNUAL_PRECIP	384071.0	2.970	30.550	32.514	14.551	
ANNUAL_TEMP_AVG	384071.0	54.833	66.467	66.207	4.201	
ANNUAL_CDD	384071.0	923.000	2612.000	2630.563	679.476	
ANNUAL_HDD	384071.0	340.000	2067.000	2158.472	908.579	
	max	skew				
STATE_ANSI	48.00	0.000				
STATE_FIPS_CODE	48.00	0.000				
ASD_CODE	99.00	0.179				
COUNTY_ANSI	507.00	0.095				
COUNTY_CODE	998.00	1.496				
...	...	...				
GROWING_SEASON_TEMP_MIN	76.05	-0.587				
ANNUAL_PRECIP	95.03	0.609				

```

ANNUAL_TEMP_AVG      77.10 -0.186
ANNUAL_CDD           4909.00  0.320
ANNUAL_HDD           4929.00  0.507

```

[96 rows x 7 columns]

Top 10 numeric columns by variance:

```

ANNUAL_HDD      825516.534321
ANNUAL_CDD      461687.084531
COUNTY_CODE     40275.811177
COUNTY_FIPS     40275.811177
HDD_JAN         35928.917592
HDD_DEC          34259.376968
HDD_FEB          32347.354784
COUNTY_ANSI      21736.955659
HDD_NOV          20741.884111
HDD_MAR          18727.997060
dtype: float64

```

Note: High variance in VALUE is expected since it measures different units (dollars, bushels, acres). Feature scaling will be important for modeling.

### 1.5.2 Outlier detection (IQR method)

```

[7]: def iqr_outliers(s):
    """Identify outliers using the IQR (Interquartile Range) method."""
    q1 = s.quantile(0.25)
    q3 = s.quantile(0.75)
    iqr = q3 - q1
    return s[(s < q1 - 1.5 * iqr) | (s > q3 + 1.5 * iqr)]


# Check key numeric columns for outliers
qa_numeric = [
    "ANNUAL_PRECIP",
    "ANNUAL_TEMP_AVG",
    "GROWING_SEASON_PRECIP",
    "VALUE_numeric",
]
qa_numeric = [c for c in qa_numeric if c in df.columns]

print("Outlier Analysis (using 1.5 x IQR rule):")
print("=" * 60)
for col in qa_numeric:
    s = df[col].dropna()
    out = iqr_outliers(s)

```

```

print(f"\n{col}: {len(out)} outliers ({len(out)/len(df):.3%} of rows)")
if len(out) > 0:
    print("Top 5 outlier values:")
    display(out.sort_values(ascending=False).head(5))

# Additional quality checks
dup_count = df.duplicated().sum()
print(f"\n'*60}")
print(f"Duplicate rows: {dup_count}")

if "ANNUAL_PRECIP" in df.columns:
    neg_precip = len(df[df["ANNUAL_PRECIP"] < 0])
    print(f"Negative precipitation rows: {neg_precip}")

print("\nNote: Climate outliers likely represent extreme weather events")
print("(droughts, floods) and should be retained for modeling.")

```

Outlier Analysis (using 1.5 x IQR rule):

---

ANNUAL\_PRECIP: 1613 outliers (0.405% of rows)

Top 5 outlier values:

1532	95.03
1097	95.03
295479	95.03
293780	95.03
390391	95.03

Name: ANNUAL\_PRECIP, dtype: float64

ANNUAL\_TEMP\_AVG: 5276 outliers (1.325% of rows)

Top 5 outlier values:

244212	77.1
1117	77.1
1466	77.1
2562	77.1
3264	77.1

Name: ANNUAL\_TEMP\_AVG, dtype: float64

GROWING\_SEASON\_PRECIP: 7381 outliers (1.854% of rows)

Top 5 outlier values:

394392	71.66
337	71.66
198	71.66
394708	71.66
385289	71.66

```
Name: GROWING_SEASON_PRECIP, dtype: float64
```

```
=====
Duplicate rows: 0
Negative precipitation rows: 0
```

```
Note: Climate outliers likely represent extreme weather events
(droughts, floods) and should be retained for modeling.
```

## 1.6 What's in the data?

```
[8]: # Quick look at the categorical variables
print(f"Counties covered: {df['COUNTY_NAME'].nunique()}")
print(f"Different crops: {df['COMMODITY_DESC'].nunique()}")
print(f"Statistic types: {df['STATISTICCAT_DESC'].nunique()}")

print(f"\nMost common crops:")
for crop, count in df["COMMODITY_DESC"].value_counts().head(5).items():
    print(f"  {crop}: {count:,} records")

print(f"\nTypes of measurements:")
for stat, count in df["STATISTICCAT_DESC"].value_counts().head(5).items():
    pct = count / len(df) * 100
    print(f"  {stat}: {count:,} ({pct:.1f}%)")
```

```
Counties covered: 256
```

```
Different crops: 165
```

```
Statistic types: 16
```

```
Most common crops:
```

```
  COTTON: 47,859 records
```

```
  WHEAT: 45,658 records
```

```
  SORGHUM: 27,754 records
```

```
  HAY: 22,679 records
```

```
  CORN: 21,495 records
```

```
Types of measurements:
```

```
  AREA HARVESTED: 166,318 (41.8%)
```

```
  PRODUCTION: 45,786 (11.5%)
```

```
  SALES: 32,474 (8.2%)
```

```
  AREA BEARING & NON-BEARING: 31,155 (7.8%)
```

```
  YIELD: 23,473 (5.9%)
```

## 1.7 Data types and climate features

```
[21]: # Distribution of data types
print("Data types:")
for dtype, count in df.dtypes.value_counts().items():
    print(f" {str(dtype):15s}: {count:3d} columns")

# Summary stats for our engineered climate features
print("\n\nClimate feature ranges:")
climate_stats = df[engineered_columns].describe().T
climate_stats["range"] = climate_stats["max"] - climate_stats["min"]
climate_stats[["min", "max", "mean", "std"]].round(2)
```

Data types:

```
float64      : 91 columns
object       : 24 columns
int64        :   6 columns
```

Climate feature ranges:

```
[21]:          min     max     mean     std
GROWING_SEASON_PRECIP    1.53    71.66   18.95    8.94
GROWING_SEASON_TEMP_AVG  67.15   85.70   77.32    3.31
GROWING_SEASON_TEMP_MAX  81.30   98.68   88.92    3.11
GROWING_SEASON_TEMP_MIN  52.68   76.05   65.70    4.26
ANNUAL_PRECIP            2.97   95.03   32.51   14.55
ANNUAL_TEMP_AVG          54.83   77.10   66.21    4.20
ANNUAL_CDD               923.00  4909.00  2630.56  679.48
ANNUAL_HDD               340.00  4929.00  2158.47  908.58
```

## 1.8 Dataset overview

```
[10]: # Put together a comprehensive summary
print("=" * 70)
print("DATASET SUMMARY".center(70))
print("=" * 70)

print(f"\nSize:")
print(f" {df.shape[0]} records x {df.shape[1]} features")
print(f" Memory: {df.memory_usage(deep=True).sum() / (1024**2):.1f} MB")

print(f"\nFeature breakdown:")
print(f" Agricultural: {len([c for c in usda_columns if c in df.columns])}")
print(f" Location: {len([c for c in location_columns if c in df.columns])}")
print(f" Time: {len([c for c in temporal_columns if c in df.columns])}")
monthly_climate = len([
    [
```

```

    c
    for c in precip_columns
    + tmax_columns
    + tmin_columns
    + tavg_columns
    + cdd_columns
    + hdd_columns
    if c in df.columns
]
)
print(f" Monthly climate: {monthly_climate}")
print(
    f" Engineered climate: {len([c for c in engineered_columns if c in df.
columns])}"
)
if "YEAR" in df.columns:
    print(f"\nTime coverage:")
    print(f" {df['YEAR'].min()}-{df['YEAR'].max()} ({df['YEAR'].nunique()} years)")

# Check census year pattern
year_counts = df["YEAR"].value_counts().sort_index()
census_years = [2002, 2007, 2012, 2017, 2022]
census_avg = year_counts[year_counts.index.isin(census_years)].mean()
non_census_avg = year_counts[~year_counts.index.isin(census_years)].mean()
print(f" Census years avg: {census_avg:.0f} records")
print(f" Other years avg: {non_census_avg:.0f} records")

if "COUNTY_NAME" in df.columns:
    print(f"\nGeography:")
    print(f" {df['COUNTY_NAME'].nunique()} counties")

if "COMMODITY_DESC" in df.columns:
    print(f"\nAgriculture:")
    print(f" {df['COMMODITY_DESC'].nunique()} commodities tracked")

if "STATISTICCAT_DESC" in df.columns:
    print(f" {df['STATISTICCAT_DESC'].nunique()} measurement types")

print(f"\nData quality:")
total_missing = df.isnull().sum().sum()
total_cells = df.shape[0] * df.shape[1]
print(f" {100 - (total_missing/total_cells)*100:.1f}% complete")
print(f" {len(missing_data)}/{len(df.columns)} columns have missing values")

print("=" * 70)

```

=====

DATASET SUMMARY

=====

Size:

398,204 records × 120 features

Memory: 842.9 MB

Feature breakdown:

Agricultural: 13

Location: 18

Time: 7

Monthly climate: 72

Engineered climate: 8

Time coverage:

2000-2023 (24 years)

Census years avg: 62,236 records

Other years avg: 4,580 records

Geography:

256 counties

Agriculture:

165 commodities tracked

16 measurement types

Data quality:

Memory: 842.9 MB

Feature breakdown:

Agricultural: 13

Location: 18

Time: 7

Monthly climate: 72

Engineered climate: 8

Time coverage:

2000-2023 (24 years)

Census years avg: 62,236 records

Other years avg: 4,580 records

Geography:

256 counties

Agriculture:

165 commodities tracked

16 measurement types


```
Data quality:  
93.0% complete  
88/120 columns have missing values
```

```
=====  
93.0% complete  
88/120 columns have missing values
```

## 1.9 Visualizations

Let's look at what crops dominate the dataset:

```
[11]: # Top commodities in the dataset  
plt.figure(figsize=(12, 8))  
df[["COMMODITY_DESC"]].value_counts().head(20).plot(kind="barh", u  
↳color="forestgreen")  
plt.title("Top 20 Commodities by Record Count", fontsize=14, fontweight="bold")  
plt.xlabel("Number of Records")  
plt.ylabel("Commodity")  
plt.grid(axis="x", alpha=0.3)  
plt.tight_layout()  
plt.show()
```



### 1.9.1 Geographic climate patterns


```
[12]: # Scatter: precipitation vs temperature by county
county_climate = (
    df.groupby("COUNTY_NAME")[["ANNUAL_PRECIP", "ANNUAL_TEMP_AVG"]].mean().
    dropna()
)

plt.figure(figsize=(12, 8))
plt.scatter(
    county_climate["ANNUAL_TEMP_AVG"],
    county_climate["ANNUAL_PRECIP"],
    alpha=0.6,
    s=100,
    c=county_climate["ANNUAL_PRECIP"],
    cmap="Blues",
    edgecolors="black",
)
plt.colorbar(label="Annual Precipitation (inches)")
plt.title(
    "County Climate Patterns: Temperature vs Precipitation",
    fontsize=14,
    fontweight="bold",
)
plt.xlabel("Annual Average Temperature (°F)")
plt.ylabel("Annual Precipitation (inches)")
plt.grid(True, alpha=0.3)

# Annotate a few extremes
wettest = county_climate.nlargest(1, "ANNUAL_PRECIP").index[0]
driest = county_climate.nsmallest(1, "ANNUAL_PRECIP").index[0]
hottest = county_climate.nlargest(1, "ANNUAL_TEMP_AVG").index[0]
coolest = county_climate.nsmallest(1, "ANNUAL_TEMP_AVG").index[0]

for county in [wettest, driest, hottest, coolest]:
    row = county_climate.loc[county]
    plt.annotate(
        county,
        (row["ANNUAL_TEMP_AVG"], row["ANNUAL_PRECIP"]),
        xytext=(5, 5),
        textcoords="offset points",
        fontsize=9,
        bbox=dict(boxstyle="round,pad=0.3", facecolor="yellow", alpha=0.7),
    )

plt.tight_layout()
plt.show()
```



### 1.9.2 What types of measurements do we have?

```
[13]: # Breakdown of statistic types
stat_counts = df["STATISTICCAT_DESC"].value_counts()

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))

# Pie chart
colors = plt.cm.Set3(range(len(stat_counts)))
ax1.pie(
    stat_counts.values,
    labels=stat_counts.index,
    autopct="%1.1f%%",
    startangle=90,
    colors=colors,
)
ax1.set_title("Distribution of Measurement Types", fontsize=14, fontweight="bold")


# Bar chart for top 10
stat_counts.head(10).plot(kind="barh", ax=ax2, color="teal")
```

```

ax2.set_title("Top 10 Measurement Types", fontsize=14, fontweight="bold")
ax2.set_xlabel("Number of Records")
ax2.set_ylabel("Measurement Type")
ax2.grid(axis="x", alpha=0.3)

plt.tight_layout()
plt.show()

```



### 1.9.3 Data collection over time

```

[14]: # Records per year - shows census year pattern
year_counts = df["YEAR"].value_counts().sort_index()
census_years = [2002, 2007, 2012, 2017, 2022]

plt.figure(figsize=(14, 6))
bars = plt.bar(
    year_counts.index, year_counts.values, color="steelblue", edgecolor="black"
)

# Highlight census years
for i, year in enumerate(year_counts.index):
    if year in census_years:
        bars[i].set_color("coral")
        bars[i].set_edgecolor("darkred")

plt.title("Records Per Year (Census Years in Orange)", fontsize=14, fontweight="bold")
plt.xlabel("Year")
plt.ylabel("Number of Records")
plt.xticks(year_counts.index, rotation=45)
plt.grid(axis="y", alpha=0.3)
plt.tight_layout()


```

```

plt.show()

print(
    f"Census years have {census_avg/non_census_avg:.1f}x more records than non-census years"
)

```



Census years have 13.6x more records than non-census years

#### 1.9.4 Climate trends over time

```
[15]: # Analyze climate trends across the 24-year period
trend_cols = [c for c in ["ANNUAL_TEMP_AVG", "ANNUAL_PRECIP"] if c in df.columns]
if "YEAR" in df.columns and len(trend_cols) > 0:
    annual = df.groupby("YEAR")[trend_cols].mean().dropna()

    plt.figure(figsize=(12, 5))
    for col in trend_cols:
        plt.plot(annual.index, annual[col], label=f"{col} (annual mean)", alpha=0.6)
        plt.plot(
            annual.index,
            annual[col].rolling(3, min_periods=1).mean(),
            label=f"{col} (3-year rolling avg)",
            linewidth=2,
        )


    plt.title(
        "Annual Climate Trends: Temperature & Precipitation (2000-2023)",
        fontsize=14,
```

```

        fontweight="bold",
    )
plt.xlabel("Year")
plt.ylabel("Value")
plt.legend()
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()

print("\nNote: 3-year rolling average smooths year-to-year variability")
print("to reveal longer-term trends. Look for overall warming or")
print("changes in precipitation patterns over the 24-year period.")

```



Note: 3-year rolling average smooths year-to-year variability  
to reveal longer-term trends. Look for overall warming or  
changes in precipitation patterns over the 24-year period.


### 1.9.5 How do climate features correlate?

```
[16]: # Correlation between our engineered climate features
climate_features = [col for col in engineered_columns if col in df.columns]
if len(climate_features) > 0:
    plt.figure(figsize=(10, 8))
    corr_matrix = df[climate_features].corr()
    sns.heatmap(
        corr_matrix,
        annot=True,
        fmt=".2f",
        cmap="coolwarm",
        center=0,
```

```

        square=True,
        linewidths=1,
    )
plt.title("Climate Feature Correlations", fontsize=14, fontweight="bold")
plt.tight_layout()
plt.show()

```



### 1.9.6 Climate variable distributions

```
[17]: # Distribution plots for key climate variables
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

# Growing season precip
axes[0, 0].hist(
    df["GROWING_SEASON_PRECIP"].dropna(), bins=50, color="steelblue", edgecolor="black"
)
```

```

)
axes[0, 0].set_title("Growing Season Precipitation", fontweight="bold")
axes[0, 0].set_xlabel("Precipitation (inches)")
axes[0, 0].set_ylabel("Frequency")
mean_precip = df["GROWING_SEASON_PRECIP"].mean()
axes[0, 0].axvline(
    mean_precip, color="red", linestyle="--", label=f'Mean: {mean_precip:.1f}'
)
axes[0, 0].legend()

# Growing season temp
axes[0, 1].hist(
    df["GROWING_SEASON_TEMP_AVG"].dropna(), bins=50, color="coral", □
    edgecolor="black"
)
axes[0, 1].set_title("Growing Season Avg Temperature", fontweight="bold")
axes[0, 1].set_xlabel("Temperature (°F)")
axes[0, 1].set_ylabel("Frequency")
mean_temp = df["GROWING_SEASON_TEMP_AVG"].mean()
axes[0, 1].axvline(
    mean_temp, color="red", linestyle="--", label=f'Mean: {mean_temp:.1f}°F'
)
axes[0, 1].legend()

# Annual precip
axes[1, 0].hist(df["ANNUAL_PRECIP"].dropna(), bins=50, color="green", □
    edgecolor="black")
axes[1, 0].set_title("Annual Precipitation", fontweight="bold")
axes[1, 0].set_xlabel("Precipitation (inches)")
axes[1, 0].set_ylabel("Frequency")
mean_annual_precip = df["ANNUAL_PRECIP"].mean()
axes[1, 0].axvline(
    mean_annual_precip,
    color="red",
    linestyle="--",
    label=f'Mean: {mean_annual_precip:.1f}'',
)
axes[1, 0].legend()

# Annual temp
axes[1, 1].hist(
    df["ANNUAL_TEMP_AVG"].dropna(), bins=50, color="orange", edgecolor="black"
)
axes[1, 1].set_title("Annual Avg Temperature", fontweight="bold")
axes[1, 1].set_xlabel("Temperature (°F)")
axes[1, 1].set_ylabel("Frequency")
mean_annual_temp = df["ANNUAL_TEMP_AVG"].mean()


```

```

        axes[1, 1].axvline(
            mean_annual_temp,
            color="red",
            linestyle="--",
            label=f"Mean: {mean_annual_temp:.1f}°F",
        )
        axes[1, 1].legend()

plt.tight_layout()
plt.show()

```



### 1.9.7 Regional climate patterns

```

[18]: # Average climate by county
county_climate = (
    df.groupby("COUNTY_NAME")[["ANNUAL_PRECIP", "ANNUAL_TEMP_AVG"]].mean().
    dropna()
)

print("Wettest counties (annual precipitation):")
print(county_climate.nlargest(10, "ANNUAL_PRECIP")[["ANNUAL_PRECIP"]].round(1))

```

```

print("\n\nDriest counties:")
print(county_climate.nsmallest(10, "ANNUAL_PRECIP") [["ANNUAL_PRECIP"]].round(1))

print("\n\nHottest counties (annual avg temp):")
print(county_climate.nlargest(10, "ANNUAL_TEMP_AVG") [["ANNUAL_TEMP_AVG"]].
    round(1))

print("\n\nCoolest counties:")
print(county_climate.nsmallest(10, "ANNUAL_TEMP_AVG") [["ANNUAL_TEMP_AVG"]].
    round(1))

```

Wettest counties (annual precipitation):

ANNUAL\_PRECIP

COUNTY\_NAME

ORANGE	69.1
JEFFERSON	66.9
HARDIN	64.2
NEWTON	61.6
GALVESTON	59.7
LIBERTY	59.6
CHAMBERS	58.9
JASPER	58.5
HARRIS	57.4
TYLER	57.3

Driest counties:

ANNUAL\_PRECIP

COUNTY\_NAME

EL PASO	9.6
HUDSPETH	10.8
REEVES	11.8
CULBERSON	11.9
LOVING	12.6
BREWSTER	13.1
WARD	14.0
PRESIDIO	14.1
GAINES	14.5
PECOS	14.6

Hottest counties (annual avg temp):

ANNUAL\_TEMP\_AVG

COUNTY\_NAME

HIDALGO	75.5
CAMERON	75.4
STARR	75.4

ZAPATA	75.3
WILLACY	75.0
KENEDY	74.7
BROOKS	74.3
JIM HOGG	74.1
KLEBERG	73.9
NUECES	73.5

Coolest counties:

COUNTY_NAME	ANNUAL_TEMP_AVG
DALLAM	56.0
HARTLEY	57.0
DEAF SMITH	57.6
SHERMAN	57.6
PARMER	58.0
OLDHAM	58.2
CASTRO	58.3
OCHILTREE	58.4
LIPSCOMB	58.4
MOORE	58.4

### 1.9.8 Does climate affect yields?

```
[19]: # Look at yield correlations for major crops
df["VALUE_numeric"] = pd.to_numeric(df["VALUE"], errors="coerce")

# Focus on yield stats for major crops
major_crops = ["COTTON", "WHEAT", "CORN", "SORGHUM"]
yield_data = df[
    (df["STATISTICCAT_DESC"] == "YIELD") & (df["COMMODITY_DESC"] .
    ↪isin(major_crops))
].copy()

if len(yield_data) > 0:
    print("Correlation: Growing Season Precipitation vs Crop Yield")
    print("-" * 60)
    for crop in major_crops:
        crop_data = yield_data[yield_data["COMMODITY_DESC"] == crop]
        if len(crop_data) > 100: # need enough data points
            corr = (
                crop_data[["GROWING_SEASON_PRECIP", "VALUE_numeric"]].corr() .
            ↪iloc[0, 1]
            )
            print(f"{crop:12s}: {corr:6.3f}")
```

```

print("\n\nCorrelation: Growing Season Temperature vs Crop Yield")
print("-" * 60)
for crop in major_crops:
    crop_data = yield_data[yield_data["COMMODITY_DESC"] == crop]
    if len(crop_data) > 100:
        corr = (
            crop_data[["GROWING_SEASON_TEMP_AVG", "VALUE_numeric"]]
            .corr()
            .iloc[0, 1]
        )
        print(f"{crop:12s}: {corr:6.3f}")

    print("\n(Values range -1 to 1: positive = yield increases with variable,"
          " negative = yield decreases with variable)")
else:
    print("Not enough yield data for correlation analysis")

```

Correlation: Growing Season Precipitation vs Crop Yield

---

COTTON	:	0.239
WHEAT	:	0.164
CORN	:	-0.166
SORGHUM	:	0.286

Correlation: Growing Season Temperature vs Crop Yield

---

COTTON	:	-0.102
WHEAT	:	-0.115
CORN	:	-0.704
SORGHUM	:	0.059

(Values range -1 to 1: positive = yield increases with variable,  
negative = yield decreases with variable)

## 1.10 Key Takeaways

**Dataset scope:** - ~398k records spanning 24 years (2000-2023) across all 256 Texas counties - 165 commodities tracked, but cotton, wheat, and sorghum dominate (top 5 crops = ~165k records) - Census years (every 5 years) have way more data - 62k avg vs 4.6k avg for other years

**Data quality:** - 93% complete overall - pretty good - Most missing values are in administrative fields (ZIP codes, regions, etc. are 100% missing) - Climate measurements are 96.5% complete

**Climate patterns we noticed:** - Temperature variables are highly correlated, which makes sense - Precipitation ranges from ~3 to 95 inches annually (mean: 32.5") - Clear regional differences: east Texas is wetter, south is hotter, west is drier - Correlations between climate and yields vary a lot by crop - corn really doesn't like heat ( $r = -0.70$ )

**For modeling:** - The strong temperature correlations suggest we could reduce dimensionality

there - Definitely need to account for census year effects - that's a huge data imbalance - The climate outliers are real weather events (droughts, floods), not errors - keep them - Might need separate models for different statistic types since "Area Harvested" dominates (42%)