

Exploratory Data Analysis: Texas Agriculture with Climate Data (2000-2023)

Project: Predicting Agricultural Statistics for Texas Counties Using Climate Data

Authors: Carter Dobbs, Johann Steinhoff, Jay Suh

Date: November 2025

Overview

This notebook performs exploratory data analysis on a merged dataset combining:

- USDA NASS QuickStats agricultural data for Texas counties
- NOAA nClimDiv monthly climate data

The goal is to understand the dataset structure, identify patterns, and prepare for predictive modeling of agricultural outcomes based on climate conditions.

```
In [1]: # Import necessary Libraries
import sys
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings

print("Python:", sys.version)
print("NumPy:", np.__version__)
print("Pandas:", pd.__version__)

# Configuration
warnings.filterwarnings('ignore')
```

```
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', 100)
plt.style.use('seaborn-v0_8-darkgrid')
sns.set_palette('husl')

print("Libraries imported successfully!")
```

Python: 3.13.9 (tags/v3.13.9:8183fa5, Oct 14 2025, 14:09:13) [MSC v.1944 64 bit (AMD64)]
NumPy: 2.3.1
Pandas: 2.3.2
Libraries imported successfully!

1. Load the Dataset

In [2]:

```
# Load the merged agriculture and climate dataset
data_path = "../data/texas_agriculture_with_climate_2000_2023.csv"
df = pd.read_csv(data_path, low_memory=False)

print(f"Dataset loaded successfully!")
print(f"Shape: {df.shape}")
print(f"Total cells: {df.shape[0] * df.shape[1]}")
```

Dataset loaded successfully!
Shape: (398204, 120)
Total cells: 47,784,480

2. Dataset Size and Structure

2.1 Basic Information

In [3]:

```
# Display dataset dimensions
print("*80")
print("DATASET DIMENSIONS")
print("*80")
print(f"Number of instances (records): {df.shape[0]}")
print(f"Number of attributes (columns): {df.shape[1]}")
```

```
print(f"Total data points: {df.shape[0] * df.shape[1]:,}")
print(f"Memory usage: {df.memory_usage(deep=True).sum() / (1024**2):.2f} MB")
print("*80")
```

=====

DATASET DIMENSIONS

=====

Number of instances (records): 398,204
Number of attributes (columns): 120
Total data points: 47,784,480
Memory usage: 842.87 MB

=====

```
In [4]: # Display first few rows
print("\nFirst 5 rows of the dataset:")
df.head()
```

First 5 rows of the dataset:

Out[4]:

	SOURCE_DESC	SECTOR_DESC	GROUP_DESC	COMMODITY_DESC	CLASS_DESC	PRODN_PRACTICE_DESC	UTIL_PRACTICE_DESC
--	-------------	-------------	------------	----------------	------------	---------------------	--------------------

0	CENSUS	CROPS	VEGETABLES	SQUASH	WINTER	ALL PRODUCTION PRACTICES	FRESH MARKET
1	CENSUS	CROPS	FRUIT & TREE NUTS	PECANS	ALL CLASSES	ALL PRODUCTION PRACTICES	ALL UTILIZATION PRACTICES
2	CENSUS	CROPS	FRUIT & TREE NUTS	TANGERINES	ALL CLASSES	ALL PRODUCTION PRACTICES	ALL UTILIZATION PRACTICES
3	CENSUS	CROPS	FRUIT & TREE NUTS	PECANS	ALL CLASSES	ALL PRODUCTION PRACTICES	ALL UTILIZATION PRACTICES
4	CENSUS	CROPS	FIELD CROPS	PEANUTS	ALL CLASSES	IRRIGATED	ALL UTILIZATION PRACTICES

In [5]:

```
# Display data types and non-null counts
print("\nColumn Information:")
df.info()
```

Column Information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398204 entries, 0 to 398203
Columns: 120 entries, SOURCE_DESC to ANNUAL_HDD
dtypes: float64(90), int64(6), object(24)
memory usage: 364.6+ MB

2.2 Column Names

```
In [6]: # Display all column names
print("\nAll Columns in Dataset:")
print("*80")
for i, col in enumerate(df.columns, 1):
    print(f"{i:3d}. {col}")
print("*80")
```

All Columns in Dataset:

- ```
=====
1. SOURCE_DESC
2. SECTOR_DESC
3. GROUP_DESC
4. COMMODITY_DESC
5. CLASS_DESC
6. PRODN_PRACTICE_DESC
7. UTIL_PRACTICE_DESC
8. STATISTICCAT_DESC
9. UNIT_DESC
10. SHORT_DESC
11. DOMAIN_DESC
12. DOMAINCAT_DESC
13. AGG_LEVEL_DESC
14. STATE_ANSI
15. STATE_FIPS_CODE
16. STATE_ALPHA
17. STATE_NAME
18. ASD_CODE
19. ASD_DESC
20. COUNTY_ANSI
21. COUNTY_CODE
22. COUNTY_NAME
23. REGION_DESC
24. ZIP_5
25. WATERSHED_CODE
26. WATERSHED_DESC
27. CONGR_DISTRICT_CODE
28. COUNTRY_CODE
29. COUNTRY_NAME
30. LOCATION_DESC
31. YEAR
32. FREQ_DESC
33. BEGIN_CODE
34. END_CODE
35. REFERENCE_PERIOD_DESC
36. WEEK_ENDING
37. LOAD_TIME
38. VALUE
39. CV%
40. COUNTY_FIPS
```

41. PRECIP\_JAN  
42. PRECIP\_FEB  
43. PRECIP\_MAR  
44. PRECIP\_APR  
45. PRECIP\_MAY  
46. PRECIP\_JUN  
47. PRECIP\_JUL  
48. PRECIP\_AUG  
49. PRECIP\_SEP  
50. PRECIP\_OCT  
51. PRECIP\_NOV  
52. PRECIP\_DEC  
53. TMAX\_JAN  
54. TMAX\_FEB  
55. TMAX\_MAR  
56. TMAX\_APR  
57. TMAX\_MAY  
58. TMAX\_JUN  
59. TMAX\_JUL  
60. TMAX\_AUG  
61. TMAX\_SEP  
62. TMAX\_OCT  
63. TMAX\_NOV  
64. TMAX\_DEC  
65. TMIN\_JAN  
66. TMIN\_FEB  
67. TMIN\_MAR  
68. TMIN\_APR  
69. TMIN\_MAY  
70. TMIN\_JUN  
71. TMIN\_JUL  
72. TMIN\_AUG  
73. TMIN\_SEP  
74. TMIN\_OCT  
75. TMIN\_NOV  
76. TMIN\_DEC  
77. TAVG\_JAN  
78. TAVG\_FEB  
79. TAVG\_MAR  
80. TAVG\_APR  
81. TAVG\_MAY  
82. TAVG\_JUN

83. TAVG\_JUL  
84. TAVG\_AUG  
85. TAVG\_SEP  
86. TAVG\_OCT  
87. TAVG\_NOV  
88. TAVG\_DEC  
89. CDD\_JAN  
90. CDD\_FEB  
91. CDD\_MAR  
92. CDD\_APR  
93. CDD\_MAY  
94. CDD\_JUN  
95. CDD\_JUL  
96. CDD\_AUG  
97. CDD\_SEP  
98. CDD\_OCT  
99. CDD\_NOV  
100. CDD\_DEC  
101. HDD\_JAN  
102. HDD\_FEB  
103. HDD\_MAR  
104. HDD\_APR  
105. HDD\_MAY  
106. HDD\_JUN  
107. HDD\_JUL  
108. HDD\_AUG  
109. HDD\_SEP  
110. HDD\_OCT  
111. HDD\_NOV  
112. HDD\_DEC  
113. GROWING\_SEASON\_PRECIP  
114. GROWING\_SEASON\_TEMP\_AVG  
115. GROWING\_SEASON\_TEMP\_MAX  
116. GROWING\_SEASON\_TEMP\_MIN  
117. ANNUAL\_PRECIP  
118. ANNUAL\_TEMP\_AVG  
119. ANNUAL\_CDD  
120. ANNUAL\_HDD

=====

### 3. Attribute Categories and Descriptions

The dataset contains attributes from three main sources:

1. **USDA Agricultural Data** - Crop information and statistics
2. **NOAA Climate Data** - Monthly temperature and precipitation
3. **Engineered Features** - Derived seasonal and annual aggregates

```
In [7]: # Categorize columns
usda_columns = [
 'SOURCE_DESC', 'SECTOR_DESC', 'GROUP_DESC', 'COMMODITY_DESC', 'CLASS_DESC',
 'PRODN_PRACTICE_DESC', 'UTIL_PRACTICE_DESC', 'STATISTICCAT_DESC', 'UNIT_DESC',
 'SHORT_DESC', 'DOMAIN_DESC', 'DOMAINCAT_DESC', 'AGG_LEVEL_DESC'
]

location_columns = [
 'STATE_ANSI', 'STATE_FIPS_CODE', 'STATE_ALPHA', 'STATE_NAME',
 'ASD_CODE', 'ASD_DESC', 'COUNTY_ANSI', 'COUNTY_CODE', 'COUNTY_NAME',
 'REGION_DESC', 'ZIP_5', 'WATERSHED_CODE', 'WATERSHED_DESC',
 'CONGR_DISTRICT_CODE', 'COUNTRY_CODE', 'COUNTRY_NAME', 'LOCATION_DESC',
 'COUNTY_FIPS'
]

temporal_columns = [
 'YEAR', 'FREQ_DESC', 'BEGIN_CODE', 'END_CODE', 'REFERENCE_PERIOD_DESC',
 'WEEK_ENDING', 'LOAD_TIME'
]

target_columns = ['VALUE', 'CV_%']

Climate columns
months = ['JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL', 'AUG', 'SEP', 'OCT', 'NOV', 'DEC']
precip_columns = [f'PRECIP_{month}' for month in months]
tmax_columns = [f'TMAX_{month}' for month in months]
tmin_columns = [f'TMIN_{month}' for month in months]
tavg_columns = [f'TAVG_{month}' for month in months]
cdd_columns = [f'CDD_{month}' for month in months]
hdd_columns = [f'HDD_{month}' for month in months]
```

```
engineered_columns = [
 'GROWING_SEASON_PRECIP', 'GROWING_SEASON_TEMP_AVG',
 'GROWING_SEASON_TEMP_MAX', 'GROWING_SEASON_TEMP_MIN',
 'ANNUAL_PRECIP', 'ANNUAL_TEMP_AVG', 'ANNUAL_CDD', 'ANNUAL_HDD'
]

print("\nAttribute Categories:")
print("*80")
print(f"USDA Agricultural Attributes: {len(usda_columns)}")
print(f"Location/Geographic Attributes: {len(location_columns)}")
print(f"Temporal Attributes: {len(temporal_columns)}")
print(f"Target Variables: {len(target_columns)}")
print(f"Monthly Precipitation: {len(precip_columns)}")
print(f"Monthly Max Temperature: {len(tmax_columns)}")
print(f"Monthly Min Temperature: {len(tmin_columns)}")
print(f"Monthly Avg Temperature: {len(tavg_columns)}")
print(f"Monthly Cooling Degree Days: {len(cdd_columns)}")
print(f"Monthly Heating Degree Days: {len(hdd_columns)}")
print(f"Engineered Climate Features: {len(engineered_columns)})")
print("*80")
print(f"Total: {len(df.columns)} columns")
```

Attribute Categories:

```
=====
USDA Agricultural Attributes: 13
Location/Geographic Attributes: 18
Temporal Attributes: 7
Target Variables: 2
Monthly Precipitation: 12
Monthly Max Temperature: 12
Monthly Min Temperature: 12
Monthly Avg Temperature: 12
Monthly Cooling Degree Days: 12
Monthly Heating Degree Days: 12
Engineered Climate Features: 8
=====
Total: 120 columns
```

---

## 4. Detailed Attribute Descriptions

## 4.1 USDA Agricultural Attributes

```
In [8]: # Describe USDA agricultural columns
attribute_descriptions = {
 'SOURCE_DESC': 'Data source (e.g., CENSUS, SURVEY)',
 'SECTOR_DESC': 'Agricultural sector (e.g., CROPS, ANIMALS & PRODUCTS)',
 'GROUP_DESC': 'Commodity group (e.g., FIELD CROPS, VEGETABLES, FRUIT & TREE NUTS)',
 'COMMODITY_DESC': 'Specific commodity/crop (e.g., CORN, WHEAT, COTTON)',
 'CLASS_DESC': 'Commodity class/variety',
 'PRODN_PRACTICE_DESC': 'Production practice (e.g., IRRIGATED, NON-IRRIGATED)',
 'UTIL_PRACTICE_DESC': 'Utilization practice (e.g., GRAIN, SILAGE)',
 'STATISTICCAT_DESC': 'Type of statistic (e.g., YIELD, PRODUCTION, AREA HARVESTED)',
 'UNIT_DESC': 'Unit of measurement (e.g., BU/ACRE, TONS, $)',
 'SHORT_DESC': 'Complete description of the statistic',
 'DOMAIN_DESC': 'Domain category',
 'DOMAINCAT_DESC': 'Domain category description',
 'AGG_LEVEL_DESC': 'Geographic aggregation level'
}

print("\nUSDA Agricultural Attribute Descriptions:")
print("*80")
for col in usda_columns:
 if col in df.columns:
 unique_count = df[col].nunique()
 print(f"\n{col}:")
 print(f" Description: {attribute_descriptions.get(col, 'N/A')}")
 print(f" Type: {df[col].dtype}")
 print(f" Unique values: {unique_count}")
 if unique_count <= 10:
 print(f" Values: {df[col].unique()[:10].tolist()}")
print("*80")
```

USDA Agricultural Attribute Descriptions:

---

SOURCE\_DESC:

Description: Data source (e.g., CENSUS, SURVEY)  
Type: object  
Unique values: 2  
Values: ['CENSUS', 'SURVEY']

SECTOR\_DESC:

Description: Agricultural sector (e.g., CROPS, ANIMALS & PRODUCTS)  
Type: object  
Unique values: 1  
Values: ['CROPS']

GROUP\_DESC:

Description: Commodity group (e.g., FIELD CROPS, VEGETABLES, FRUIT & TREE NUTS)  
Type: object  
Unique values: 5  
Values: ['VEGETABLES', 'FRUIT & TREE NUTS', 'FIELD CROPS', 'CROP TOTALS', 'HORTICULTURE']

COMMODITY\_DESC:

Description: Specific commodity/crop (e.g., CORN, WHEAT, COTTON)  
Type: object  
Unique values: 165

CLASS\_DESC:

Description: Commodity class/variety  
Type: object  
Unique values: 85

PRODN\_PRACTICE\_DESC:

Description: Production practice (e.g., IRRIGATED, NON-IRRIGATED)  
Type: object  
Unique values: 8  
Values: ['ALL PRODUCTION PRACTICES', 'IRRIGATED', 'NON-IRRIGATED', 'IN THE OPEN', 'UNDER PROTECTION', 'IN THE OPEN, IRRIGATED', 'PRODUCTION CONTRACT', 'ORGANIC']

UTIL\_PRACTICE\_DESC:

Description: Utilization practice (e.g., GRAIN, SILAGE)  
Type: object  
Unique values: 20

STATISTICCAT\_DESC:  
Description: Type of statistic (e.g., YIELD, PRODUCTION, AREA HARVESTED)  
Type: object  
Unique values: 16

UNIT\_DESC:  
Description: Unit of measurement (e.g., BU/ACRE, TONS, \$)  
Type: object  
Unique values: 21

SHORT\_DESC:  
Description: Complete description of the statistic  
Type: object  
Unique values: 1442

DOMAIN\_DESC:  
Description: Domain category  
Type: object  
Unique values: 7  
Values: ['TOTAL', 'AREA HARVESTED', 'AREA HARVESTED, FRESH MARKET & PROCESSING', 'NAICS CLASSIFICATION', 'AREA BEARING & NON-BEARING', 'BALES GINNED', 'ORGANIC STATUS']

DOMAINCAT\_DESC:  
Description: Domain category description  
Type: object  
Unique values: 38

AGG\_LEVEL\_DESC:  
Description: Geographic aggregation level  
Type: object  
Unique values: 1  
Values: ['COUNTY']  
=====

## 4.2 Location and Geographic Attributes

```
In [9]: # Analyze location columns
print("\nLocation/Geographic Attributes:")
print("*"*80)
location_desc = {
```

```
'STATE_NAME': 'State name',
'COUNTY_NAME': 'County name',
'COUNTY_FIPS': 'Federal Information Processing Standard county code',
'ASD_DESC': 'Agricultural Statistics District',
'REGION_DESC': 'Geographic region'
}

for col in ['STATE_NAME', 'COUNTY_NAME', 'COUNTY_FIPS', 'ASD_DESC']:
 if col in df.columns:
 print(f"\n{col}:")
 print(f" Description: {location_desc.get(col, 'Geographic identifier')}")
 print(f" Unique values: {df[col].nunique()}")
 if col == 'COUNTY_NAME':
 print(f" Sample counties: {df[col].unique()[:10].tolist()}")
print("*"*80)
```

Location/Geographic Attributes:

```
=====
```

STATE\_NAME:

Description: State name  
Unique values: 1

COUNTY\_NAME:

Description: County name  
Unique values: 256  
Sample counties: ['BELL', 'SHACKELFORD', 'KERR', 'MONTAGUE', 'DAWSON', 'OTHER (COMBINED) COUNTIES', 'HALE', 'TERRY',  
, 'WOOD', 'LIMESTONE']

COUNTY\_FIPS:

Description: Federal Information Processing Standard county code  
Unique values: 255

ASD\_DESC:

Description: Agricultural Statistics District  
Unique values: 15

## 4.3 Temporal Attributes

```
In [10]: # Analyze temporal attributes
```

```
print("\nTemporal Attributes:")
print("*80)
if 'YEAR' in df.columns:
 print(f"\nYEAR:")
 print(f" Description: Year of observation")
 print(f" Range: {df['YEAR'].min()} to {df['YEAR'].max()}")
 print(f" Unique years: {df['YEAR'].nunique()}")
 print(f" Years covered: {sorted(df['YEAR'].unique())}")
print("*80)
```

Temporal Attributes:

```
=====
```

YEAR:

```
 Description: Year of observation
 Range: 2000 to 2023
 Unique years: 24
 Years covered: [np.int64(2000), np.int64(2001), np.int64(2002), np.int64(2003), np.int64(2004), np.int64(2005), n
p.int64(2006), np.int64(2007), np.int64(2008), np.int64(2009), np.int64(2010), np.int64(2011), np.int64(2012), np.int
64(2013), np.int64(2014), np.int64(2015), np.int64(2016), np.int64(2017), np.int64(2018), np.int64(2019), np.int64(20
20), np.int64(2021), np.int64(2022), np.int64(2023)]
=====
```

## 4.4 Target Variable

```
In [11]: # Analyze target variable
print("\nTarget Variable:")
print("*80)
if 'VALUE' in df.columns:
 # Convert VALUE to numeric, handling non-numeric values
 df['VALUE_numeric'] = pd.to_numeric(df['VALUE'], errors='coerce')

 print(f"\nVALUE:")
 print(f" Description: Agricultural statistic value (varies by STATISTICCAT_DESC and UNIT_DESC)")
 print(f" Type: {df['VALUE'].dtype}")
 print(f" Non-null count: {df['VALUE'].notna().sum():,}")
 print(f" Null count: {df['VALUE'].isna().sum():,}")
 print(f" Numeric values: {df['VALUE_numeric'].notna().sum():,}")
 print(f"\n Numeric VALUE Statistics:")
 print(f" Min: {df['VALUE_numeric'].min()}")
 print(f" Max: {df['VALUE_numeric'].max()})
```

```
print(f" Mean: {df['VALUE_numeric'].mean():.2f}")
print(f" Median: {df['VALUE_numeric'].median():.2f}")
print(f" Std Dev: {df['VALUE_numeric'].std():.2f}")

Check for non-numeric values
non_numeric = df[df['VALUE_numeric'].isna() & df['VALUE'].notna()][['VALUE']].unique()
if len(non_numeric) > 0:
 print(f"\n Non-numeric values found: {non_numeric[:10]}")
print("=*80")
```

Target Variable:

---

VALUE:

Description: Agricultural statistic value (varies by STATISTICCAT\_DESC and UNIT\_DESC)  
Type: object  
Non-null count: 398,204  
Null count: 0  
Numeric values: 210,970

Numeric VALUE Statistics:

Min: 0.0  
Max: 999.0  
Mean: 68.16  
Median: 6.00  
Std Dev: 168.92

Non-numeric values found: ['(D)' '1,919' '102,851,000' '1,500' '5,500' '1,448,000' '40,192' '3,000'  
'22,000' '4,378']

---

## 4.5 Climate Attributes - Monthly Precipitation

In [12]:

```
Analyze monthly precipitation
print("\nMonthly Precipitation Attributes (inches):")
print("=*80")
precip_stats = pd.DataFrame()
for col in precip_columns:
 if col in df.columns:
 precip_stats[col] = [
 df[col].min(),
```

```
 df[col].max(),
 df[col].mean(),
 df[col].std()
]

precip_stats.index = ['Min', 'Max', 'Mean', 'Std Dev']
print(precip_stats.T.round(2))
print("*80)
```

Monthly Precipitation Attributes (inches):

|            | Min | Max   | Mean | Std Dev |
|------------|-----|-------|------|---------|
| PRECIP_JAN | 0.0 | 9.79  | 2.33 | 2.12    |
| PRECIP_FEB | 0.0 | 11.45 | 1.65 | 1.47    |
| PRECIP_MAR | 0.0 | 11.74 | 2.99 | 2.35    |
| PRECIP_APR | 0.0 | 11.88 | 2.22 | 1.67    |
| PRECIP_MAY | 0.0 | 23.41 | 3.52 | 2.36    |
| PRECIP_JUN | 0.0 | 16.20 | 3.32 | 2.48    |
| PRECIP_JUL | 0.0 | 16.30 | 3.27 | 2.95    |
| PRECIP_AUG | 0.0 | 46.04 | 3.91 | 4.71    |
| PRECIP_SEP | 0.0 | 19.74 | 2.72 | 1.99    |
| PRECIP_OCT | 0.0 | 23.18 | 2.63 | 2.51    |
| PRECIP_NOV | 0.0 | 17.76 | 1.88 | 2.04    |
| PRECIP_DEC | 0.0 | 11.91 | 2.09 | 1.98    |

## 4.6 Climate Attributes - Monthly Temperatures

```
In [13]: # Analyze temperature ranges
print("\nMonthly Temperature Attributes (Fahrenheit):")
print("*80)

Sample a few months for detailed display
sample_months = ['JAN', 'APR', 'JUL', 'OCT']
temp_summary = []

for month in sample_months:
 tmax_col = f'TMAX_{month}'
 tmin_col = f'TMIN_{month}'
 tavg_col = f'TAVG_{month}'
```

```
if all(col in df.columns for col in [tmax_col, tmin_col, tavg_col]):
 temp_summary.append({
 'Month': month,
 'TMAX_Min': df[tmax_col].min(),
 'TMAX_Max': df[tmax_col].max(),
 'TMAX_Mean': df[tmax_col].mean(),
 'TMIN_Min': df[tmin_col].min(),
 'TMIN_Max': df[tmin_col].max(),
 'TMIN_Mean': df[tmin_col].mean(),
 'TAVG_Mean': df[tavg_col].mean()
 })

temp_df = pd.DataFrame(temp_summary)
print(temp_df.round(1))
print("\nShowing sample months: January, April, July, October")
print("=*80)
```

Monthly Temperature Attributes (Fahrenheit):

```
=====
 Month TMAX_Min TMAX_Max TMAX_Mean TMIN_Min TMIN_Max TMIN_Mean \\\n0 JAN 38.2 78.8 59.8 17.1 56.8 35.5
1 APR 63.4 95.2 78.9 32.3 70.2 54.3
2 JUL 84.8 106.3 94.3 60.8 80.3 72.0
3 OCT 62.3 92.8 79.0 34.5 71.3 54.2
```

```
TAVG_Mean
0 47.6
1 66.6
2 83.1
3 66.6
```

(Showing sample months: January, April, July, October)

## 4.7 Climate Attributes - Degree Days

```
In [14]: # Analyze degree days
print("\nDegree Days Attributes:")
print("=*80")
print("\nCooling Degree Days (CDD) - Energy needed for cooling:")
cdd_stats = pd.DataFrame()
```

```
for col in ['CDD_JAN', 'CDD_APR', 'CDD_JUL', 'CDD_OCT']:
 if col in df.columns:
 cdd_stats[col] = [df[col].min(), df[col].max(), df[col].mean()]
cdd_stats.index = ['Min', 'Max', 'Mean']
print(cdd_stats.T.round(1))

print("\nHeating Degree Days (HDD) - Energy needed for heating:")
hdd_stats = pd.DataFrame()
for col in ['HDD_JAN', 'HDD_APR', 'HDD_JUL', 'HDD_OCT']:
 if col in df.columns:
 hdd_stats[col] = [df[col].min(), df[col].max(), df[col].mean()]
hdd_stats.index = ['Min', 'Max', 'Mean']
print(hdd_stats.T.round(1))
print("=*80")
```

Degree Days Attributes:

---

Cooling Degree Days (CDD) - Energy needed for cooling:

|         | Min   | Max   | Mean  |
|---------|-------|-------|-------|
| CDD_JAN | 0.0   | 174.0 | 13.2  |
| CDD_APR | 0.0   | 491.0 | 127.8 |
| CDD_JUL | 269.0 | 864.0 | 561.6 |
| CDD_OCT | 0.0   | 492.0 | 128.2 |

Heating Degree Days (HDD) - Energy needed for heating:

|         | Min  | Max    | Mean  |
|---------|------|--------|-------|
| HDD_JAN | 99.0 | 1148.0 | 551.6 |
| HDD_APR | 0.0  | 476.0  | 79.7  |
| HDD_JUL | 0.0  | 0.0    | 0.0   |
| HDD_OCT | 0.0  | 477.0  | 78.1  |

---

## 4.8 Engineered Climate Features

```
In [15]: # Analyze engineered features
print("\nEngineered Climate Features:")
print("=*80")
eng_desc = {
 'GROWING_SEASON_PRECIP': 'Total precipitation April-September (inches)',
 'GROWING_SEASON_TEMP_AVG': 'Average temperature April-September (°F)',
```

```
'GROWING_SEASON_TEMP_MAX': 'Average max temperature April-September (°F)',
'GROWING_SEASON_TEMP_MIN': 'Average min temperature April-September (°F)',
'ANNUAL_PRECIP': 'Total annual precipitation (inches)',
'ANNUAL_TEMP_AVG': 'Annual average temperature (°F)',
'ANNUAL_CDD': 'Total annual cooling degree days',
'ANNUAL_HDD': 'Total annual heating degree days'
}

for col in engineered_columns:
 if col in df.columns:
 print(f"\n{col}:")
 print(f" Description: {eng_desc.get(col, 'N/A')}")
 print(f" Range: {df[col].min():.2f} to {df[col].max():.2f}")
 print(f" Mean: {df[col].mean():.2f}")
 print(f" Std Dev: {df[col].std():.2f}")
print("*"*80)
```

Engineered Climate Features:

=====

GROWING\_SEASON\_PRECIP:

Description: Total precipitation April-September (inches)  
Range: 1.53 to 71.66  
Mean: 18.95  
Std Dev: 8.94

GROWING\_SEASON\_TEMP\_AVG:

Description: Average temperature April-September (°F)  
Range: 67.15 to 85.70  
Mean: 77.32  
Std Dev: 3.31

GROWING\_SEASON\_TEMP\_MAX:

Description: Average max temperature April-September (°F)  
Range: 81.30 to 98.68  
Mean: 88.92  
Std Dev: 3.11

GROWING\_SEASON\_TEMP\_MIN:

Description: Average min temperature April-September (°F)  
Range: 52.68 to 76.05  
Mean: 65.70  
Std Dev: 4.26

ANNUAL\_PRECIP:

Description: Total annual precipitation (inches)  
Range: 2.97 to 95.03  
Mean: 32.51  
Std Dev: 14.55

ANNUAL\_TEMP\_AVG:

Description: Annual average temperature (°F)  
Range: 54.83 to 77.10  
Mean: 66.21  
Std Dev: 4.20

ANNUAL\_CDD:

Description: Total annual cooling degree days  
Range: 923.00 to 4909.00

Mean: 2630.56  
Std Dev: 679.48

ANNUAL\_HDD:  
Description: Total annual heating degree days  
Range: 340.00 to 4929.00  
Mean: 2158.47  
Std Dev: 908.58

---

## 5. Missing Values Analysis

```
In [16]: # Check for missing values
missing_data = pd.DataFrame({
 'Column': df.columns,
 'Missing_Count': df.isnull().sum().values,
 'Missing_Percentage': (df.isnull().sum().values / len(df) * 100)
})
missing_data = missing_data[missing_data['Missing_Count'] > 0].sort_values(
 'Missing_Percentage', ascending=False
)

print("\nMissing Values Summary:")
print("*80")
if len(missing_data) > 0:
 print(f"\nColumns with missing values: {len(missing_data)}")
 print(missing_data.to_string(index=False))
else:
 print("No missing values found!")
print("*80")
```

## Missing Values Summary:

=====

| Column              | Missing_Count | Missing_Percentage |
|---------------------|---------------|--------------------|
| REGION_DESC         | 398204        | 100.000000         |
| WEEK_ENDING         | 398204        | 100.000000         |
| ZIP_5               | 398204        | 100.000000         |
| WATERSHED_DESC      | 398204        | 100.000000         |
| CONGR_DISTRICT_CODE | 398204        | 100.000000         |
| CV_%                | 211970        | 53.231509          |
| VALUE_numeric       | 187234        | 47.019618          |
| PRECIP_JAN          | 14133         | 3.549186           |
| COUNTY_ANSI         | 14133         | 3.549186           |
| PRECIP_OCT          | 14133         | 3.549186           |
| PRECIP_MAR          | 14133         | 3.549186           |
| PRECIP_APR          | 14133         | 3.549186           |
| PRECIP_MAY          | 14133         | 3.549186           |
| PRECIP_JUN          | 14133         | 3.549186           |
| PRECIP_JUL          | 14133         | 3.549186           |
| PRECIP_AUG          | 14133         | 3.549186           |
| PRECIP_SEP          | 14133         | 3.549186           |
| TMAX_FEB            | 14133         | 3.549186           |
| PRECIP_NOV          | 14133         | 3.549186           |
| PRECIP_DEC          | 14133         | 3.549186           |
| TMAX_JAN            | 14133         | 3.549186           |
| TMAX_APR            | 14133         | 3.549186           |
| TMAX_MAR            | 14133         | 3.549186           |
| TMAX_MAY            | 14133         | 3.549186           |
| PRECIP_FEB          | 14133         | 3.549186           |
| CDD_JAN             | 14133         | 3.549186           |
| TMAX_JUN            | 14133         | 3.549186           |
| TMAX_AUG            | 14133         | 3.549186           |
| TMAX_JUL            | 14133         | 3.549186           |
| TMAX_OCT            | 14133         | 3.549186           |
| TMAX_NOV            | 14133         | 3.549186           |
| TMAX_DEC            | 14133         | 3.549186           |
| TMAX_SEP            | 14133         | 3.549186           |
| TMIN_FEB            | 14133         | 3.549186           |
| TMIN_MAR            | 14133         | 3.549186           |
| TMIN_APR            | 14133         | 3.549186           |
| TMIN_MAY            | 14133         | 3.549186           |

|                       |       |          |
|-----------------------|-------|----------|
| TMIN_JUN              | 14133 | 3.549186 |
| TMIN_JUL              | 14133 | 3.549186 |
| TMIN_AUG              | 14133 | 3.549186 |
| TMIN_JAN              | 14133 | 3.549186 |
| TMIN_OCT              | 14133 | 3.549186 |
| TMIN_NOV              | 14133 | 3.549186 |
| TMIN_DEC              | 14133 | 3.549186 |
| TAVG_JAN              | 14133 | 3.549186 |
| TAVG_FEB              | 14133 | 3.549186 |
| TAVG_MAR              | 14133 | 3.549186 |
| TAVG_APR              | 14133 | 3.549186 |
| TAVG_MAY              | 14133 | 3.549186 |
| TAVG_JUN              | 14133 | 3.549186 |
| TAVG_JUL              | 14133 | 3.549186 |
| TAVG_AUG              | 14133 | 3.549186 |
| TAVG_SEP              | 14133 | 3.549186 |
| TAVG_OCT              | 14133 | 3.549186 |
| TAVG_NOV              | 14133 | 3.549186 |
| TAVG_DEC              | 14133 | 3.549186 |
| TMIN_SEP              | 14133 | 3.549186 |
| HDD_MAY               | 14133 | 3.549186 |
| CDD_FEB               | 14133 | 3.549186 |
| CDD_APR               | 14133 | 3.549186 |
| CDD_MAR               | 14133 | 3.549186 |
| CDD_JUN               | 14133 | 3.549186 |
| CDD_JUL               | 14133 | 3.549186 |
| CDD_AUG               | 14133 | 3.549186 |
| CDD_MAY               | 14133 | 3.549186 |
| CDD_OCT               | 14133 | 3.549186 |
| CDD_NOV               | 14133 | 3.549186 |
| CDD_DEC               | 14133 | 3.549186 |
| HDD_JAN               | 14133 | 3.549186 |
| HDD_FEB               | 14133 | 3.549186 |
| HDD_MAR               | 14133 | 3.549186 |
| HDD_APR               | 14133 | 3.549186 |
| CDD_SEP               | 14133 | 3.549186 |
| GROWING_SEASON_PRECIP | 14133 | 3.549186 |
| HDD_JUN               | 14133 | 3.549186 |
| HDD_AUG               | 14133 | 3.549186 |
| HDD_JUL               | 14133 | 3.549186 |
| HDD_OCT               | 14133 | 3.549186 |
| HDD_NOV               | 14133 | 3.549186 |

|                         |       |          |
|-------------------------|-------|----------|
| HDD_DEC                 | 14133 | 3.549186 |
| HDD_SEP                 | 14133 | 3.549186 |
| ANNUAL_PRECIP           | 14133 | 3.549186 |
| GROWING_SEASON_TEMP_AVG | 14133 | 3.549186 |
| GROWING_SEASON_TEMP_MIN | 14133 | 3.549186 |
| GROWING_SEASON_TEMP_MAX | 14133 | 3.549186 |
| ANNUAL_CDD              | 14133 | 3.549186 |
| ANNUAL_TEMP_AVG         | 14133 | 3.549186 |
| ANNUAL_HDD              | 14133 | 3.549186 |
| ASD_DESC                | 261   | 0.065544 |
| =====                   |       |          |

## 6. Key Categorical Variables

```
In [17]: # Analyze key categorical variables
print("\nKey Categorical Variables:")
print("*"*80)

Counties
if 'COUNTY_NAME' in df.columns:
 print(f"\nUnique Counties: {df['COUNTY_NAME'].nunique()}")
 print(f"Top 10 counties by record count:")
 print(df['COUNTY_NAME'].value_counts().head(10))

Crops
if 'COMMODITY_DESC' in df.columns:
 print(f"\n\nUnique Commodities/Crops: {df['COMMODITY_DESC'].nunique()}")
 print(f"Top 10 commodities by record count:")
 print(df['COMMODITY_DESC'].value_counts().head(10))

Statistics Types
if 'STATISTICCAT_DESC' in df.columns:
 print(f"\n\nUnique Statistic Types: {df['STATISTICCAT_DESC'].nunique()}")
 print(f"Statistic types:")
 print(df['STATISTICCAT_DESC'].value_counts())

print("*"*80)
```

**Key Categorical Variables:**

Unique Counties: 256

Top 10 counties by record count:

| COUNTY_NAME               |       |
|---------------------------|-------|
| OTHER (COMBINED) COUNTIES | 13872 |
| HIDALGO                   | 3424  |
| LUBBOCK                   | 3012  |
| MEDINA                    | 2961  |
| BRAZORIA                  | 2937  |
| CAMERON                   | 2892  |
| WHARTON                   | 2824  |
| WILLIAMSON                | 2783  |
| HUNT                      | 2776  |
| COLLIN                    | 2763  |

Name: count, dtype: int64

Unique Commodities/Crops: 165

Top 10 commodities by record count:

| COMMODITY_DESC   |       |
|------------------|-------|
| COTTON           | 47859 |
| WHEAT            | 45658 |
| SORGHUM          | 27754 |
| HAY              | 22679 |
| CORN             | 21495 |
| PECANS           | 14528 |
| VEGETABLE TOTALS | 14198 |
| HAY & HAYLAGE    | 12663 |
| HAYLAGE          | 11293 |
| OATS             | 8002  |

Name: count, dtype: int64

Unique Statistic Types: 16

Statistic types:

| STATISTICCAT_DESC          |        |
|----------------------------|--------|
| AREA HARVESTED             | 166318 |
| PRODUCTION                 | 45786  |
| SALES                      | 32474  |
| AREA BEARING & NON-BEARING | 31155  |

```
YIELD 23473
AREA PLANTED 23340
AREA BEARING 20622
AREA NON-BEARING 18228
AREA IN PRODUCTION 16082
GINNED BALES 10992
AREA GROWN 4342
CAPACITY 3452
ACTIVE GINS 871
AREA NOT HARVESTED 518
OPERATIONS 494
SUCROSE 57
Name: count, dtype: int64
=====
```

---

## 7. Data Summary Statistics

```
In [18]: # Comprehensive numerical summary
print("\nNumerical Attributes Summary:")
print("*"*80)
numeric_cols = df.select_dtypes(include=[np.number]).columns.tolist()
print(f"Total numerical columns: {len(numeric_cols)}")
df[numeric_cols].describe().T
```

```
Numerical Attributes Summary:
=====
```

```
Total numerical columns: 97
```

Out[18]:

|  |                            | count    | mean         | std        | min          | 25%          | 50%          | 75%          |
|--|----------------------------|----------|--------------|------------|--------------|--------------|--------------|--------------|
|  | <b>STATE_ANSI</b>          | 398204.0 | 48.000000    | 0.000000   | 48.000000    | 48.000000    | 48.000000    | 48.000000    |
|  | <b>STATE_FIPS_CODE</b>     | 398204.0 | 48.000000    | 0.000000   | 48.000000    | 48.000000    | 48.000000    | 48.000000    |
|  | <b>ASD_CODE</b>            | 398204.0 | 49.134444    | 27.868408  | 11.000000    | 22.000000    | 51.000000    | 81.000000    |
|  | <b>COUNTY_ANSI</b>         | 384071.0 | 246.935694   | 147.434581 | 1.000000     | 121.000000   | 233.000000   | 371.000000   |
|  | <b>COUNTY_CODE</b>         | 398204.0 | 273.592362   | 200.688343 | 1.000000     | 125.000000   | 245.000000   | 387.000000   |
|  | <b>REGION_DESC</b>         | 0.0      | NaN          | NaN        | NaN          | NaN          | NaN          | NaN          |
|  | <b>ZIP_5</b>               | 0.0      | NaN          | NaN        | NaN          | NaN          | NaN          | NaN          |
|  | <b>WATERSHED_CODE</b>      | 398204.0 | 0.000000     | 0.000000   | 0.000000     | 0.000000     | 0.000000     | 0.000000     |
|  | <b>WATERSHED_DESC</b>      | 0.0      | NaN          | NaN        | NaN          | NaN          | NaN          | NaN          |
|  | <b>CONGR_DISTRICT_CODE</b> | 0.0      | NaN          | NaN        | NaN          | NaN          | NaN          | NaN          |
|  | <b>COUNTRY_CODE</b>        | 398204.0 | 9000.000000  | 0.000000   | 9000.000000  | 9000.000000  | 9000.000000  | 9000.000000  |
|  | <b>YEAR</b>                | 398204.0 | 2011.601332  | 7.110870   | 2000.000000  | 2007.000000  | 2012.000000  | 2017.000000  |
|  | <b>BEGIN_CODE</b>          | 398204.0 | 0.336506     | 1.748010   | 0.000000     | 0.000000     | 0.000000     | 0.000000     |
|  | <b>END_CODE</b>            | 398204.0 | 0.336506     | 1.748010   | 0.000000     | 0.000000     | 0.000000     | 0.000000     |
|  | <b>WEEK_ENDING</b>         | 0.0      | NaN          | NaN        | NaN          | NaN          | NaN          | NaN          |
|  | <b>COUNTY_FIPS</b>         | 398204.0 | 48273.592362 | 200.688343 | 48001.000000 | 48125.000000 | 48245.000000 | 48387.000000 |
|  | <b>PRECIP_JAN</b>          | 384071.0 | 2.329363     | 2.118643   | 0.000000     | 0.630000     | 1.680000     | 3.490000     |
|  | <b>PRECIP_FEB</b>          | 384071.0 | 1.648177     | 1.469168   | 0.000000     | 0.490000     | 1.290000     | 2.450000     |
|  | <b>PRECIP_MAR</b>          | 384071.0 | 2.986804     | 2.349424   | 0.000000     | 1.080000     | 2.460000     | 4.570000     |
|  | <b>PRECIP_APR</b>          | 384071.0 | 2.219726     | 1.669245   | 0.000000     | 0.940000     | 1.890000     | 3.130000     |
|  | <b>PRECIP_MAY</b>          | 384071.0 | 3.518074     | 2.355157   | 0.000000     | 1.850000     | 2.990000     | 4.750000     |
|  | <b>PRECIP_JUN</b>          | 384071.0 | 3.323892     | 2.480152   | 0.000000     | 1.530000     | 2.780000     | 4.520000     |

|                   | <b>count</b> | <b>mean</b> | <b>std</b> | <b>min</b> | <b>25%</b> | <b>50%</b> | <b>75%</b> |
|-------------------|--------------|-------------|------------|------------|------------|------------|------------|
| <b>PRECIP_JUL</b> | 384071.0     | 3.269744    | 2.954237   | 0.000000   | 1.060000   | 2.410000   | 4.520000   |
| <b>PRECIP_AUG</b> | 384071.0     | 3.907348    | 4.712599   | 0.000000   | 1.570000   | 2.740000   | 4.500000   |
| <b>PRECIP_SEP</b> | 384071.0     | 2.715322    | 1.993747   | 0.000000   | 1.360000   | 2.210000   | 3.640000   |
| <b>PRECIP_OCT</b> | 384071.0     | 2.628202    | 2.509809   | 0.000000   | 0.850000   | 1.980000   | 3.410000   |
| <b>PRECIP_NOV</b> | 384071.0     | 1.879319    | 2.044262   | 0.000000   | 0.410000   | 1.120000   | 2.680000   |
| <b>PRECIP_DEC</b> | 384071.0     | 2.088385    | 1.982861   | 0.000000   | 0.500000   | 1.450000   | 3.290000   |
| <b>TMAX_JAN</b>   | 384071.0     | 59.790140   | 6.592459   | 38.200000  | 55.800000  | 59.900000  | 64.200000  |
| <b>TMAX_FEB</b>   | 384071.0     | 62.866700   | 6.584740   | 38.100000  | 58.300000  | 62.300000  | 67.100000  |
| <b>TMAX_MAR</b>   | 384071.0     | 72.339708   | 4.797473   | 54.800000  | 69.500000  | 73.000000  | 75.500000  |
| <b>TMAX_APR</b>   | 384071.0     | 78.900251   | 4.902214   | 63.400000  | 75.700000  | 79.100000  | 82.200000  |
| <b>TMAX_MAY</b>   | 384071.0     | 85.506729   | 4.202680   | 72.400000  | 82.800000  | 85.500000  | 87.900000  |
| <b>TMAX_JUN</b>   | 384071.0     | 91.953301   | 3.665832   | 81.100000  | 89.100000  | 92.000000  | 94.600000  |
| <b>TMAX_JUL</b>   | 384071.0     | 94.268269   | 4.319146   | 84.800000  | 91.000000  | 94.000000  | 97.600000  |
| <b>TMAX_AUG</b>   | 384071.0     | 94.033791   | 3.132952   | 83.000000  | 92.000000  | 94.200000  | 96.100000  |
| <b>TMAX_SEP</b>   | 384071.0     | 88.879373   | 2.940199   | 75.100000  | 87.500000  | 89.100000  | 90.500000  |
| <b>TMAX_OCT</b>   | 384071.0     | 79.047023   | 4.735964   | 62.300000  | 76.000000  | 79.900000  | 82.300000  |
| <b>TMAX_NOV</b>   | 384071.0     | 68.938123   | 5.565167   | 50.600000  | 65.100000  | 69.400000  | 72.800000  |
| <b>TMAX_DEC</b>   | 384071.0     | 61.027564   | 6.086587   | 40.200000  | 57.200000  | 60.700000  | 65.200000  |
| <b>TMIN_JAN</b>   | 384071.0     | 35.462742   | 7.411738   | 17.100000  | 30.500000  | 35.100000  | 40.600000  |
| <b>TMIN_FEB</b>   | 384071.0     | 38.059913   | 8.482063   | 15.200000  | 32.300000  | 37.100000  | 43.900000  |
| <b>TMIN_MAR</b>   | 384071.0     | 47.390848   | 8.024373   | 23.700000  | 41.700000  | 48.300000  | 53.400000  |
| <b>TMIN_APR</b>   | 384071.0     | 54.288962   | 6.743861   | 32.300000  | 49.800000  | 55.200000  | 59.000000  |

|                 | <b>count</b> | <b>mean</b> | <b>std</b> | <b>min</b> | <b>25%</b> | <b>50%</b> | <b>75%</b> |
|-----------------|--------------|-------------|------------|------------|------------|------------|------------|
| <b>TMIN_MAY</b> | 384071.0     | 62.119878   | 5.517217   | 44.600000  | 58.900000  | 63.000000  | 65.800000  |
| <b>TMIN_JUN</b> | 384071.0     | 69.238010   | 3.787790   | 55.300000  | 67.500000  | 69.700000  | 71.700000  |
| <b>TMIN_JUL</b> | 384071.0     | 71.953154   | 3.459330   | 60.800000  | 70.100000  | 72.400000  | 74.400000  |
| <b>TMIN_AUG</b> | 384071.0     | 71.527746   | 3.599886   | 57.600000  | 69.900000  | 72.400000  | 73.900000  |
| <b>TMIN_SEP</b> | 384071.0     | 65.079088   | 4.285552   | 49.400000  | 62.900000  | 65.500000  | 67.800000  |
| <b>TMIN_OCT</b> | 384071.0     | 54.174450   | 5.439506   | 34.500000  | 51.100000  | 54.000000  | 57.300000  |
| <b>TMIN_NOV</b> | 384071.0     | 44.794034   | 6.946607   | 24.200000  | 40.200000  | 45.000000  | 49.700000  |
| <b>TMIN_DEC</b> | 384071.0     | 37.213173   | 7.175075   | 16.700000  | 32.500000  | 37.600000  | 42.000000  |
| <b>TAVG_JAN</b> | 384071.0     | 47.630627   | 6.673176   | 28.000000  | 43.000000  | 47.500000  | 51.900000  |
| <b>TAVG_FEB</b> | 384071.0     | 50.467742   | 7.295234   | 28.300000  | 45.400000  | 49.700000  | 55.500000  |
| <b>TAVG_MAR</b> | 384071.0     | 59.868461   | 6.116887   | 42.200000  | 55.500000  | 60.800000  | 64.000000  |
| <b>TAVG_APR</b> | 384071.0     | 66.598578   | 5.440260   | 49.100000  | 62.800000  | 67.100000  | 70.300000  |
| <b>TAVG_MAY</b> | 384071.0     | 73.818140   | 4.352670   | 59.000000  | 71.100000  | 74.100000  | 76.600000  |
| <b>TAVG_JUN</b> | 384071.0     | 80.601499   | 3.182502   | 68.300000  | 78.600000  | 80.400000  | 82.800000  |
| <b>TAVG_JUL</b> | 384071.0     | 83.114065   | 3.383755   | 73.600000  | 80.600000  | 83.100000  | 85.500000  |
| <b>TAVG_AUG</b> | 384071.0     | 82.783886   | 2.927166   | 71.000000  | 81.100000  | 83.300000  | 84.700000  |
| <b>TAVG_SEP</b> | 384071.0     | 76.984091   | 3.334562   | 62.200000  | 75.500000  | 77.400000  | 79.000000  |
| <b>TAVG_OCT</b> | 384071.0     | 66.614999   | 4.658557   | 49.600000  | 63.900000  | 66.900000  | 69.600000  |
| <b>TAVG_NOV</b> | 384071.0     | 56.871136   | 5.922910   | 37.400000  | 52.900000  | 57.300000  | 60.700000  |
| <b>TAVG_DEC</b> | 384071.0     | 49.126560   | 6.395236   | 29.600000  | 45.300000  | 49.100000  | 53.450000  |
| <b>CDD_JAN</b>  | 384071.0     | 13.182521   | 22.239548  | 0.000000   | 0.000000   | 6.000000   | 16.000000  |
| <b>CDD_FEB</b>  | 384071.0     | 16.954032   | 31.980921  | 0.000000   | 0.000000   | 5.000000   | 19.000000  |

|                | <b>count</b> | <b>mean</b> | <b>std</b> | <b>min</b> | <b>25%</b> | <b>50%</b> | <b>75%</b> |             |
|----------------|--------------|-------------|------------|------------|------------|------------|------------|-------------|
| <b>CDD_MAR</b> | 384071.0     | 70.807023   | 59.400854  | 0.000000   | 21.000000  | 60.000000  | 104.000000 | 111.000000  |
| <b>CDD_APR</b> | 384071.0     | 127.774258  | 90.050001  | 0.000000   | 55.000000  | 115.000000 | 179.000000 | 211.000000  |
| <b>CDD_MAY</b> | 384071.0     | 286.313281  | 116.900177 | 17.000000  | 202.000000 | 284.000000 | 359.000000 | 431.000000  |
| <b>CDD_JUN</b> | 384071.0     | 468.594637  | 94.147183  | 139.000000 | 408.000000 | 463.000000 | 535.000000 | 607.000000  |
| <b>CDD_JUL</b> | 384071.0     | 561.560446  | 104.890538 | 269.000000 | 484.000000 | 560.000000 | 636.000000 | 711.000000  |
| <b>CDD_AUG</b> | 384071.0     | 551.406982  | 90.440251  | 198.000000 | 500.000000 | 568.000000 | 611.000000 | 671.000000  |
| <b>CDD_SEP</b> | 384071.0     | 363.454101  | 93.565537  | 39.000000  | 318.000000 | 374.000000 | 421.000000 | 471.000000  |
| <b>CDD_OCT</b> | 384071.0     | 128.241148  | 83.068407  | 0.000000   | 71.000000  | 120.000000 | 171.000000 | 211.000000  |
| <b>CDD_NOV</b> | 384071.0     | 33.587509   | 41.358282  | 0.000000   | 6.000000   | 20.000000  | 45.000000  | 71.000000   |
| <b>CDD_DEC</b> | 384071.0     | 8.686683    | 19.340051  | 0.000000   | 0.000000   | 0.000000   | 10.000000  | 19.000000   |
| <b>HDD_JAN</b> | 384071.0     | 551.612514  | 189.549248 | 99.000000  | 418.000000 | 547.000000 | 680.000000 | 1111.000000 |
| <b>HDD_FEB</b> | 384071.0     | 424.013893  | 179.853704 | 21.000000  | 284.000000 | 431.000000 | 549.000000 | 1111.000000 |
| <b>HDD_MAR</b> | 384071.0     | 229.891205  | 136.850272 | 11.000000  | 133.000000 | 191.000000 | 315.000000 | 411.000000  |
| <b>HDD_APR</b> | 384071.0     | 79.723429   | 83.083749  | 0.000000   | 19.000000  | 50.000000  | 119.000000 | 171.000000  |
| <b>HDD_MAY</b> | 384071.0     | 12.491977   | 24.019026  | 0.000000   | 0.000000   | 0.000000   | 14.000000  | 17.000000   |
| <b>HDD_JUN</b> | 384071.0     | 0.452885    | 2.471190   | 0.000000   | 0.000000   | 0.000000   | 0.000000   | 0.000000    |
| <b>HDD_JUL</b> | 384071.0     | 0.000000    | 0.000000   | 0.000000   | 0.000000   | 0.000000   | 0.000000   | 0.000000    |
| <b>HDD_AUG</b> | 384071.0     | 0.027347    | 0.461584   | 0.000000   | 0.000000   | 0.000000   | 0.000000   | 0.000000    |
| <b>HDD_SEP</b> | 384071.0     | 3.543553    | 8.293623   | 0.000000   | 0.000000   | 0.000000   | 5.000000   | 11.000000   |
| <b>HDD_OCT</b> | 384071.0     | 78.098724   | 68.371846  | 0.000000   | 29.000000  | 60.000000  | 104.000000 | 141.000000  |
| <b>HDD_NOV</b> | 384071.0     | 277.654038  | 144.020429 | 12.000000  | 172.000000 | 249.000000 | 370.000000 | 481.000000  |
| <b>HDD_DEC</b> | 384071.0     | 500.962536  | 185.092887 | 27.000000  | 368.000000 | 493.000000 | 612.000000 | 1111.000000 |

|                                | count    | mean        | std        | min        | 25%         | 50%         | 75%         |
|--------------------------------|----------|-------------|------------|------------|-------------|-------------|-------------|
| <b>GROWING_SEASON_PRECIP</b>   | 384071.0 | 18.954105   | 8.938067   | 1.530000   | 12.660000   | 17.020000   | 23.690000   |
| <b>GROWING_SEASON_TEMP_AVG</b> | 384071.0 | 77.316710   | 3.309409   | 67.150000  | 75.483333   | 77.600000   | 79.583333   |
| <b>GROWING_SEASON_TEMP_MAX</b> | 384071.0 | 88.923619   | 3.107366   | 81.300000  | 86.633333   | 88.833333   | 91.100000   |
| <b>GROWING_SEASON_TEMP_MIN</b> | 384071.0 | 65.701140   | 4.260499   | 52.683333  | 63.650000   | 66.466667   | 68.400000   |
| <b>ANNUAL_PRECIP</b>           | 384071.0 | 32.514355   | 14.550911  | 2.970000   | 20.560000   | 30.550000   | 43.280000   |
| <b>ANNUAL_TEMP_AVG</b>         | 384071.0 | 66.206649   | 4.201270   | 54.833333  | 63.800000   | 66.466667   | 68.841667   |
| <b>ANNUAL_CDD</b>              | 384071.0 | 2630.562620 | 679.475595 | 923.000000 | 2205.000000 | 2612.000000 | 3040.000000 |
| <b>ANNUAL_HDD</b>              | 384071.0 | 2158.472100 | 908.579405 | 340.000000 | 1488.000000 | 2067.000000 | 2684.000000 |
| <b>VALUE_numeric</b>           | 210970.0 | 68.155548   | 168.917532 | 0.000000   | 2.000000    | 6.000000    | 32.000000   |

## 8. Dataset Summary Report

### Complete Dataset Overview

In [19]:

```
Generate comprehensive summary report
print("\n" + "="*80)
print("COMPREHENSIVE DATASET SUMMARY REPORT")
print("="*80)
print(f"\n1. DATASET SIZE")
print(f" - Total instances: {df.shape[0]}:")
print(f" - Total attributes: {df.shape[1]}:")
print(f" - Total data points: {df.shape[0] * df.shape[1]}")
print(f" - Memory usage: {df.memory_usage(deep=True).sum() / (1024**2):.2f} MB")

print(f"\n2. ATTRIBUTE BREAKDOWN")
print(f" - USDA Agricultural: {len([c for c in usda_columns if c in df.columns])}")
print(f" - Location/Geographic: {len([c for c in location_columns if c in df.columns])}")
print(f" - Temporal: {len([c for c in temporal_columns if c in df.columns])}")
print(f" - Climate (Monthly): {len([c for c in precip_columns + tmax_columns + tmin_columns + tavg_columns + cdd_cc])}
```

```
print(f" - Engineered Features: {len([c for c in engineered_columns if c in df.columns])}")
print(f" - Target Variables: {len([c for c in target_columns if c in df.columns])}")

if 'YEAR' in df.columns:
 print(f"\n3. TEMPORAL COVERAGE")
 print(f" - Years: {df['YEAR'].min()} to {df['YEAR'].max()}")
 print(f" - Total years: {df['YEAR'].nunique()}")

if 'COUNTY_NAME' in df.columns:
 print(f"\n4. GEOGRAPHIC COVERAGE")
 print(f" - Unique counties: {df['COUNTY_NAME'].nunique()}")

if 'COMMODITY_DESC' in df.columns:
 print(f"\n5. AGRICULTURAL COVERAGE")
 print(f" - Unique commodities/crops: {df['COMMODITY_DESC'].nunique()}")

if 'STATISTICCAT_DESC' in df.columns:
 print(f" - Statistic types: {df['STATISTICCAT_DESC'].nunique()}")

print("\n6. DATA QUALITY")
total_missing = df.isnull().sum().sum()
total_cells = df.shape[0] * df.shape[1]
print(f" - Total missing values: {total_missing:,}")
print(f" - Missing percentage: {((total_missing/total_cells)*100:.2f}%)")
print(f" - Columns with missing values: {len(missing_data)}")

print("\n" + "*80)
print("END OF SUMMARY REPORT")
print("*80)
```

=====

COMPREHENSIVE DATASET SUMMARY REPORT

=====

1. DATASET SIZE

- Total instances: 398,204
- Total attributes: 121
- Total data points: 48,182,684
- Memory usage: 845.91 MB

2. ATTRIBUTE BREAKDOWN

- USDA Agricultural: 13
- Location/Geographic: 18
- Temporal: 7
- Climate (Monthly): 72
- Engineered Features: 8
- Target Variables: 2

3. TEMPORAL COVERAGE

- Years: 2000 to 2023
- Total years: 24

4. GEOGRAPHIC COVERAGE

- Unique counties: 256

5. AGRICULTURAL COVERAGE

- Unique commodities/crops: 165
- Statistic types: 16

6. DATA QUALITY

- Total missing values: 3,535,258
- Missing percentage: 7.34%
- Columns with missing values: 89

=====

END OF SUMMARY REPORT

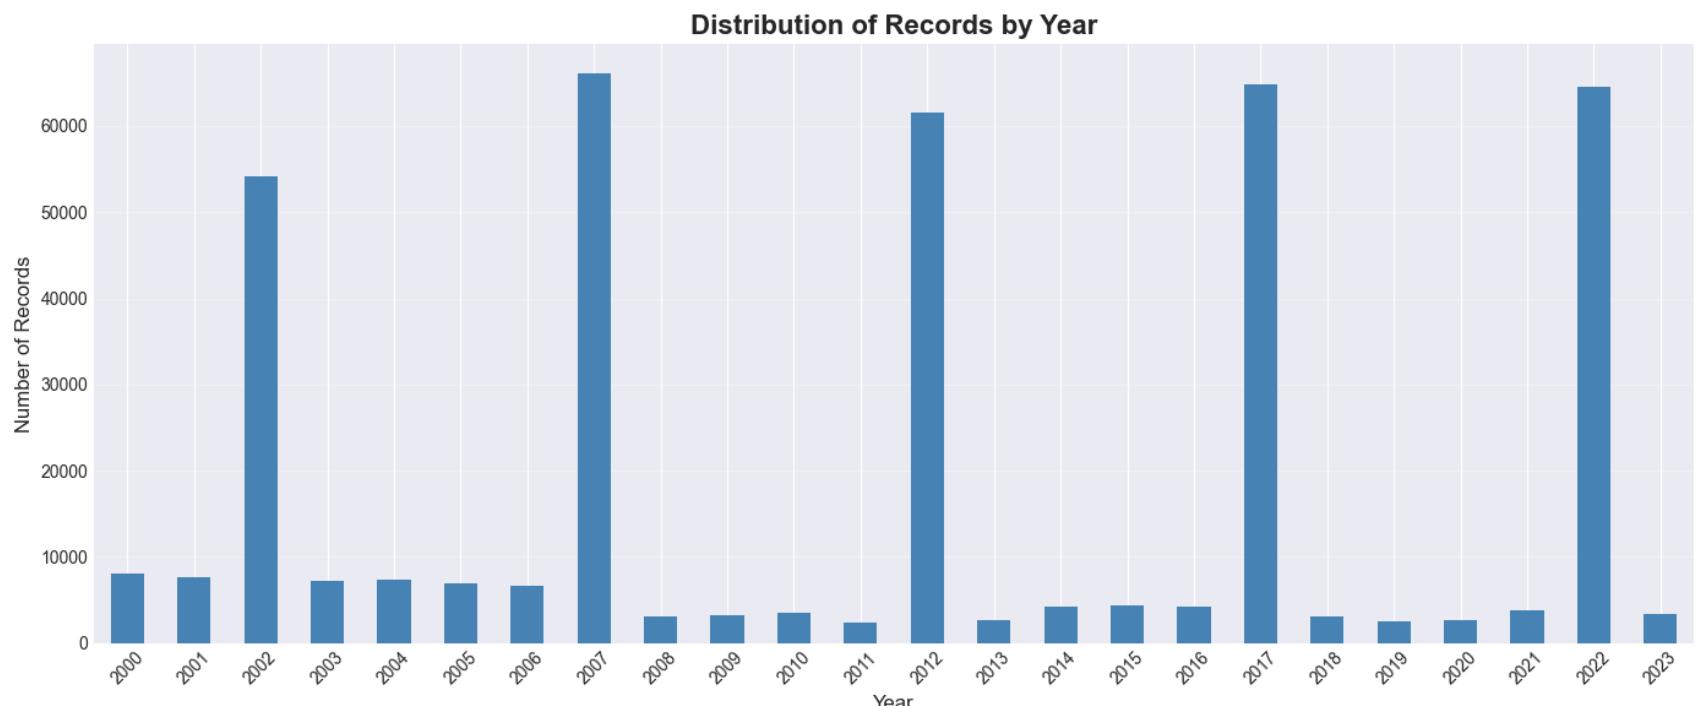
=====

---

## 9. Data Visualization (Optional)

## 9.1 Distribution of Records by Year

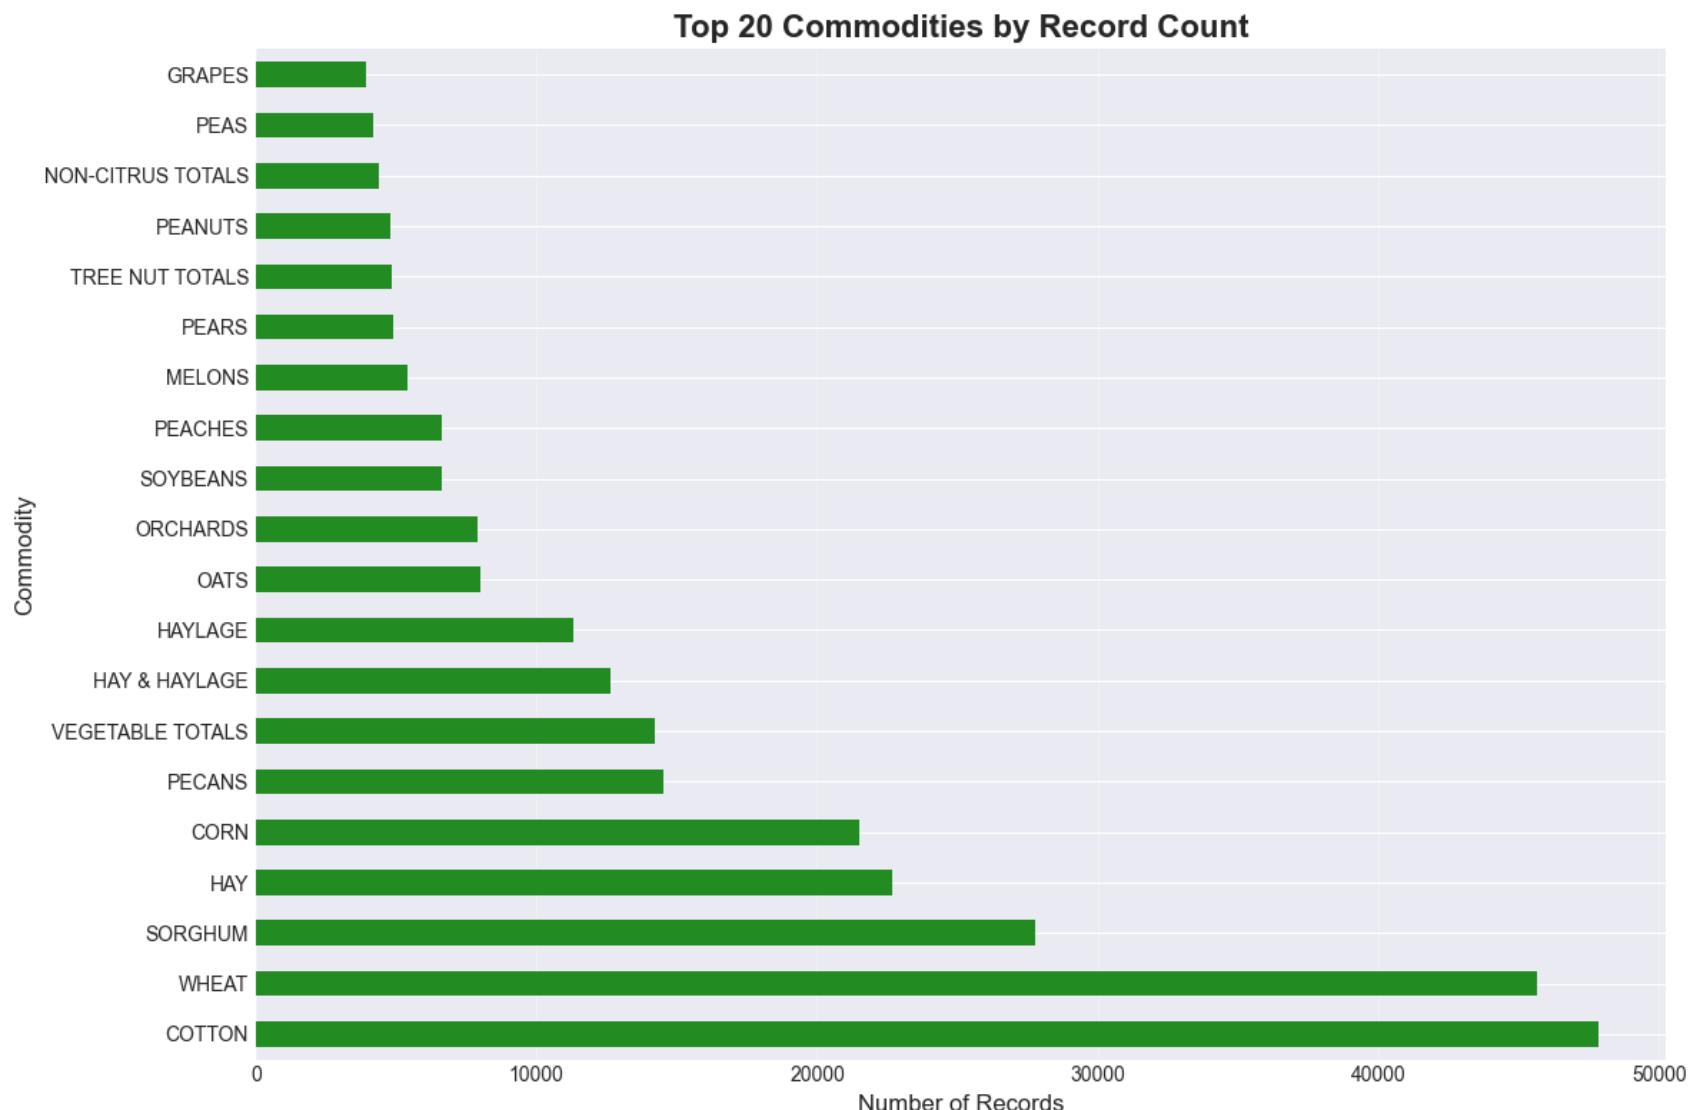
```
In [20]: # Plot records by year
if 'YEAR' in df.columns:
 plt.figure(figsize=(14, 6))
 df['YEAR'].value_counts().sort_index().plot(kind='bar', color='steelblue')
 plt.title('Distribution of Records by Year', fontsize=16, fontweight='bold')
 plt.xlabel('Year', fontsize=12)
 plt.ylabel('Number of Records', fontsize=12)
 plt.xticks(rotation=45)
 plt.grid(axis='y', alpha=0.3)
 plt.tight_layout()
 plt.show()
```



## 9.2 Top Commodities

```
In [21]: # Plot top commodities
```

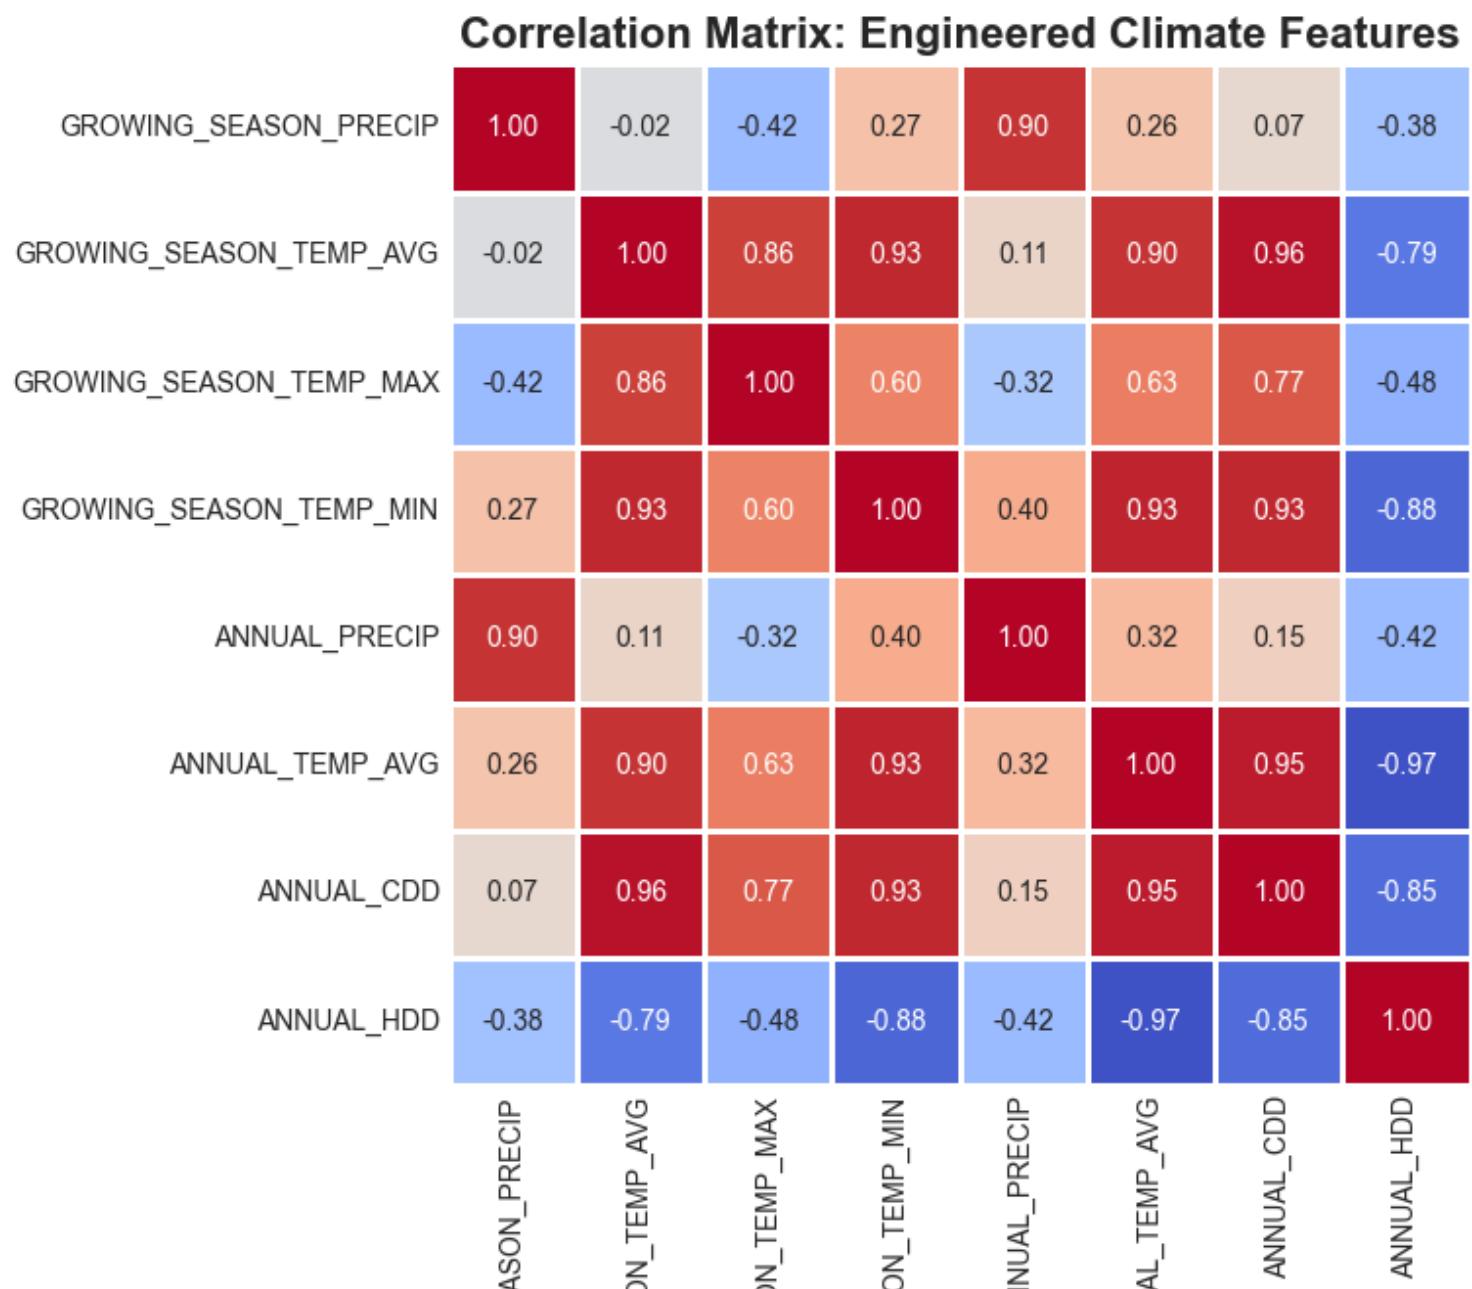
```
if 'COMMODITY_DESC' in df.columns:
 plt.figure(figsize=(12, 8))
 df['COMMODITY_DESC'].value_counts().head(20).plot(kind='barh', color='forestgreen')
 plt.title('Top 20 Commodities by Record Count', fontsize=16, fontweight='bold')
 plt.xlabel('Number of Records', fontsize=12)
 plt.ylabel('Commodity', fontsize=12)
 plt.grid(axis='x', alpha=0.3)
 plt.tight_layout()
 plt.show()
```



### 9.3 Climate Feature Correlation Heatmap

```
In [22]: # Plot correlation heatmap for engineered features
climate_features = [col for col in engineered_columns if col in df.columns]
if len(climate_features) > 0:
 plt.figure(figsize=(10, 8))
 correlation_matrix = df[climate_features].corr()
```

```
sns.heatmap(correlation_matrix, annot=True, fmt='.2f', cmap='coolwarm',
 center=0, square=True, linewidths=1)
plt.title('Correlation Matrix: Engineered Climate Features', fontsize=16, fontweight='bold')
plt.tight_layout()
plt.show()
```





---

## 10. Conclusions

### Key Findings:

1. **Dataset Size:** The dataset contains approximately 398,000 records with 120+ attributes, well exceeding the 10M data point requirement (47M+ data points).
2. **Temporal Coverage:** Data spans from 2000-2023 (or as shown in analysis above), providing substantial temporal variation for time-series analysis.
3. **Geographic Coverage:** Covers 255 Texas counties, enabling county-level predictions.
4. **Agricultural Diversity:** Contains 165 different crop types with 16 different statistical measures (yield, production, acres harvested, etc.).
5. **Climate Variables:**
  - 72 monthly climate measurements (precipitation, temperature, degree days)
  - 8 engineered seasonal/annual aggregates
  - Complete coverage of growing season (April-September) metrics
6. **Data Quality:** Analysis of missing values and data types shows areas requiring preprocessing.

### Next Steps:

1. Handle missing values and non-numeric VALUE entries
2. Encode categorical variables
3. Feature selection and engineering
4. Implement baseline models (Decision Tree Regression)
5. Apply AdaBoost for improvement

## 6. Conduct PCA for dimensionality reduction